

Internal Repository Application Implementation Guide

Purpose and outcomes

This internal guide is for maintainers, reviewers, and operators of this repository application. It documents how the application implements an auditable set of AI Credit (AIC) budget and organization policy profiles.

The fields, rates, heuristics, pages, and generated artifacts described here are implementation details. They are not released customer requirements. Customers implementing the governance approach without this application should use the [standalone customer implementation guide](#).

The application process runs from source-data validation through approved profile design, pilot, rollout, monitoring, and controlled change.

The implementation should produce these outcomes:

- One approved enterprise spending limit
- One Universal User-Level Budget (Universal ULB) baseline
- Approved cost-center ULB overrides for eligible user cohorts
- Approved organization policy profiles mapped to business requirements
- A profile register with named owners, approvers, assignments, and evidence
- Reconciled assessment, simulator, and billing records
- Measurable alert, escalation, exception, rollback, and review controls

Important: The application is an assessment, simulation, recommendation, and reporting tool. It does not create, update, or delete GitHub budgets, cost centers, organizations, teams, users, or Copilot policy settings. Authorized customer administrators perform and verify all production changes in approved systems.

Application architecture and workflow

Internal AI Credits governance implementation

■ evidence
 ■ control
 ■ approval
 ■ monitoring

Repository application route: evidence -> calculation -> approval -> manual change -> controlled operation

1

Scope and read-only evidence intake

Define the assessment

Scope, period, populations, owners and approved evidence location.

Read permitted sources

Usage, licenses, governance records, organizations, teams, and users.

Preserve provenance

Permissions, scope, retrieval time, warnings, and unknown values.

2

Validate source authority and completeness

Decision gate

Is the evidence complete and authorized for this decision?

Continue

Reconcile the approved scope and pass evidence to the route.

No: resolve or limit

Retain Unknown, assign an owner, or block approval until resolved.

3

Run the application workflow and preserve rate basis

Application route

Assessment -> Simulator -> Insights -> Recommendations -> Report

Labeled outputs

Pool, burn, exhaustion, overage, trends, and advisory signals.

Reconcile basis

Assessment promotional rates vs. simulator standard rates. Record both.

4

Separate financial controls from access-policy proposals

Financial control proposals

Enterprise spending limit
Universal ULB + higher-usage cohort overrides

Access-policy proposals

Organization profiles are planning examples, not budgets or current provider policy state.

Versioned proposal register: evidence, owners, assignments, approvals, review dates, and rate basis.

5

Reconcile, approve, and preserve the hard boundary

Reconcile before approval

Sources, rate basis, calculations, assignments, and variances.

Owners approve?

Yes: authorize change.
No: revise or defer.

Manual boundary

An authorized administrator configures the provider.

6

Verify, pilot, and stage the rollout

Independent verification

Scope, precedence, alerts, enforcement.

Limited pilot

Approved success and rollback criteria.

Stage and operate

Verify each wave and register evidence.

7

Monitor, control exceptions, and reassess

Usage, cost exposure, alerts, assignments, and drift

Exception? Roll back or contain.

Reconcile, reassess, and redesign.

Application role: assess, simulate, recommend, and report. It never mutates provider governance resources.

Diagram [IG-INT-01](#) summarizes the application architecture and implementation workflow. The linked SVG is the maintained diagram asset.

Canonical tables remain authoritative for exact fields and values, and the application does not perform provider mutations.

Data and action boundaries

Treat every value according to its provenance. Do not present calculated, synthesized, or manually entered values as raw GitHub telemetry.

Category	Values and activities	Required treatment
GitHub-sourced data	AI Credit usage by organization and model, identifiable user usage when returned, license inventory, budgets, cost centers, organizations, teams, users, and team membership	Record retrieval time, scope, warnings, and permissions. Preserve unavailable values as unknown, not zero.
Application calculations	Included pool estimate, burn rate, exhaustion day, overage credits and cost, population totals, and governance impact	Record input values, rate basis, calculation date, and scenario name. Reconcile results before approval.
Advisory heuristics	Concentration score, top-user cohorts, organization profile suggestions, governance gaps, readiness score, and daily trend	Label as advisory or synthesized. Require human review before production assignment or financial decisions.
Manual customer actions	Define ownership, approve values, configure GitHub resources and policies, communicate changes, train users, and roll back	Execute through customer change control. Capture operator, approval, timestamp, evidence, and verification.

The assessment reads GitHub governance and inventory data where credentials and endpoints permit. It does not retrieve or apply Copilot policy settings. Policy fields in Classes 7 through 10 are built-in planning examples, not a representation of current GitHub policy state.

The assessment daily trend evenly distributes current total consumption across the selected period. It is synthesized for display and is not daily GitHub telemetry. The governance readiness score is also heuristic: recommendations reduce a starting score of 100 according to their priority.

Prerequisites and accountability

Complete the prerequisites before collecting data or selecting profiles.

Prerequisite	Acceptance condition
Authorized application access	GitHub OAuth is connected by a user who can read the target enterprise and organizations.
Enterprise billing access	An approved enterprise billing credential is available when budgets, cost centers, enterprise usage, and enterprise-direct license data are required.
Defined assessment scope	Enterprise slug, organization slugs, billing cycle, assessment period, and included or excluded populations are documented.
Named business ownership	Budget owner, security owner, license owner, business owners, and platform administrator are assigned.
Approved evidence location	A controlled location exists for reports, approvals, profile-register versions, and change records.
Current naming standard	Customer naming patterns and approved values are published and available to profile designers.

Prerequisite	Acceptance condition
Change and rollback process	Production operators can submit, approve, verify, and reverse changes through the customer process.

Accountable roles

Role	Accountability
Executive sponsor	Accepts enterprise risk posture and unresolved business exceptions
Enterprise budget owner	Approves spending limits, ULB values, alert recipients, and financial variance
GitHub enterprise administrator	Performs and verifies approved GitHub configuration changes
Security and compliance owner	Approves model, feature, integration, data-handling, and exception requirements
License owner	Validates license inventory, SKU assignment, and entitlement assumptions
Organization business owner	Confirms user need, organization assignment, and operational impact
Cost-center owner	Approves cohort membership, higher ULB ceilings, and chargeback mapping
Service owner	Owns the profile register, evidence, operating cadence, and change control
Pilot lead	Coordinates pilot users, support, feedback, success measures, and rollback
Internal audit or risk reviewer	Reviews evidence completeness, approvals, reconciliations, and control operation when required

No advisory score can replace business, security, license, and budget owner approval. Record each approval before production assignment.

Required inputs and validation

Assessment inputs

Input	Implemented check	Customer control
Enterprise slug	Required; alphanumeric characters and internal hyphens; no leading or trailing hyphen; up to 39 characters	Confirm it is the enterprise account slug, not a display name or organization slug.
Organization slugs	At least one; comma-separated in the UI; each value follows the enterprise-slug character rule	Confirm completeness against the approved in-scope organization list.
Enterprise billing token	Optional in the form; 512 characters or fewer	Use an approved secret path. Never place the token in reports, registers, screenshots, or source files.
Assessment period	7, 30, or custom; custom period from 1 through 90 days	Record that usage endpoints request the current billing month; do not infer historic daily telemetry from the display period.

Simulator inputs

Input	Implemented check	Customer control
Enterprise name	Required; 100 characters or fewer	Use the approved reporting name.
License counts	Nonnegative integers; at least one licensed user across Business, Enterprise, Cloud Agent, and Spark	Reconcile counts to the license-owner record. Simulator counts are manually confirmed planning assumptions.
Billing-cycle start date	Valid date	Match the financial reporting cycle used for reconciliation.

Input	Implemented check	Customer control
Current day of cycle	Integer from 1 through 31	Confirm the day aligns with the consumed-credit snapshot.
Credits consumed so far	Nonnegative number	Enter a sourced snapshot. Mark unavailable input as unknown outside the simulator rather than entering zero.
Population allocation	Nonnegative integers whose total equals total licensed users	Assign every licensed user once across Universal, Overage, Abundant, and Exponential cohorts.
Scenario name	60 characters or fewer	Include rate basis, cycle, and scenario purpose in the supporting record.

Warning: From June 1 through August 31, 2026, the live assessment estimates included credits at 3,000 per Copilot Business license and 7,000 per Copilot Enterprise license. The simulator and generated PDF use standard rates of 1,900 Business, 3,900 Enterprise, 3,900 Cloud Agent, and 500 Spark credits per license. Label the rate basis on every comparison. Review and rerun all affected assessments, scenarios, approvals, and reports after August 31, 2026. Do not approve a variance until rate-basis differences are isolated.

Three-tier budget hierarchy

ULB means User-Level Budget. AIC means AI Credit. The implemented governance hierarchy applies controls in this order:

Tier	Control	Scope	Purpose
1	Enterprise Spending Limit	Metered overage across the enterprise	Provides the final enterprise hard cap after the included pool is consumed.
2	Universal ULB	Each user across the enterprise	Provides the default monthly per-user ceiling of 5,000 credits.
3	Cost Center ULB Overrides	Approved users in mapped cost centers	Replaces the Universal ULB with a 6,000, 7,000, or 8,000 monthly per-user ceiling.

The Tier 3 assignment is the most specific per-user control and takes precedence over the Tier 2 default for mapped users. Tier 1 remains the enterprise backstop across all metered overage. A user reaching an applicable ULB hard stop cannot continue metered usage under that profile. If aggregate metered overage reaches the enterprise spending hard stop, affected usage stops regardless of remaining user-level headroom. Validate exact provider behavior in the target environment before rollout.

Class 1 represents the included pool. It is not a configurable hard stop in the built-in catalog, but its exhaustion determines when metered overage exposure begins.

Canonical profile catalog

The ten built-in classes are canonical application records. Internal field names and string values are shown exactly where they are implemented.

Budget and pool Classes 1 through 6

Class	name	slug	scope	budgetType	skus
1	Enterprise Included AI Credits Pool	enterprise-included-ai-credits-pool	enterprise	included	All AI Credit SKUs
2	Enterprise Spending Limit	enterprise-spending-limit	enterprise	metered-coverage	All AI Credit SKUs
3	Universal User-Level Budget (ULB)	universal-ulb	enterprise	ulb	All AI Credit SKUs
4	Overage Users (ULB Cost Center)	ulb-cost-center-coverage-users	cost-center	ulb	All AI Credit SKUs
5	Abundant Users (ULB Cost Center)	ulb-cost-center-abundant-users	cost-center	ulb	All AI Credit SKUs
6	Exponential Users (ULB Cost Center)	ulb-cost-center-exponential-users	cost-center	ulb	All AI Credit SKUs

Class	budgetCredits	unit	frequency	stopUsage	excludeCostCenter	alertThresholds
1	Not specified	Not specified	Not specified	not-applicable	not-applicable	0.75 , 0.9
2	Not specified	Not specified	Not specified	true	Not specified	0.75 , 0.9
3	5000	per-user	per-month	true	Not specified	0.75 , 0.9
4	6000	per-user	per-month	true	Not specified	0.75 , 0.9
5	7000	per-user	per-month	true	Not specified	0.75 , 0.9
6	8000	per-user	per-month	true	Not specified	0.75 , 0.9

Class	costCenterId	enterpriseTeam	modelOption1	modelOption2	modelOption3
1	Not specified	Not specified	Not specified	Not specified	Not specified
2	Not specified	Not specified	Not specified	Not specified	Not specified
3	Not specified	Not specified	auto-mode	Not specified	Not specified
4	aic-0011-ovr	aic-coverage-usage-team	auto-mode	open-ai-gpt-5.6-sol	open-ai-gpt-5.6-terra
5	aic-0012-abd	aic-abundant-usage-team	auto-mode	open-ai-gpt-5.6-sol	open-ai-5.6-luna
6	aic-0013-exp	aic-exponential-usage-team	auto-mode	open-ai-gpt-5.6-terra	open-ai-5.6-luna

The 5,000, 6,000, 7,000, and 8,000 values, mappings, teams, and model strings are fixed built-in examples. Class 2 intentionally has no canonical limit; the customer selects and approves that value. Class 1 is calculated from license inputs and the applicable rate basis. Production owners may propose different values, but must version them as customer profiles and preserve the canonical source record for comparison.

Organization policy Classes 7 through 10

Classes 7 through 10 are organization policy profiles. They are never organization budget classes. Their license counts, included-credit values, cost centers, organizations, teams, models, and features are fixed examples in the

application catalog. Customers select and approve production mappings and settings after validating requirements and provider capabilities.

Class	name	slug	scope	budgetType	skus
7	Organization Policy: Standard (Copilot Business)	org-policy-standard	organization	org-policy	GHCP Business
8	Organization Policy: Engineering (Copilot Enterprise)	org-policy-engineering	organization	org-policy	GHCP Enterprise
9	Organization Policy: Architects (Copilot Enterprise)	org-policy-architects	organization	org-policy	GHCP Enterprise
10	Organization Policy: AI Platform SMEs (Copilot Enterprise, All Access)	org-policy-ai-platform	organization	org-policy	GHCP Enterprise

Class	licenseCountBusiness	licenseCountEnterprise	includedCredits	costCenterId	enterpriseOrg	enterpriseTeam
7	10	Not specified	19000	org-0001-dig	autocloudarc-digital-services	digital-services-team
8	Not specified	3	5700	org-0004-mfg	autocloudarc-manufacturing-group	manufacturing-group-team
9	Not specified	5	19500	org-0009-spc	autocloudarc-space-fleet	space-fleet-team
10	Not specified	3	11700	org-0005-mar	autocloudarc-marine-industries	marine-industries-team

Class	chatEnabled	autoModeEnabled	agentModeEnabled	cloudAgentEnabled	repoRefactorEnabled	mcpIntegrationsEnabled
7	true	true	false	false	false	false
8	true	true	limited	false	true	Not specified
9	true	true	true	limited	true	true
10	true	true	true	true	true	true

Class	modelOption1	modelOption2	modelOption3
7	gpt-5-mini	claude-haiku	Not specified
8	gpt-5	gpt-5-codex	claude-sonnet
9	gpt-5	claude-sonnet	Not specified
10	all	all	all

Fields shown as **Not specified** are undefined in the implementation. Do not convert them to **false**, zero, an empty policy, or an inferred value.

Advisory assignment rules

Use engine output to identify candidates, then apply business context and approval. Engine thresholds are advisory heuristics, not production policy.

Class	Implemented advisory signal	Required human decision
1	Included pool is estimated from license inventory and configured rates.	License and budget owners reconcile the pool to authoritative billing records.
2	Recommend when projected pool exhaustion is before day 25 of a 30-day cycle; suggested limit is 120% of projected overage cost.	Budget owner selects the limit and confirms cash exposure, continuity needs, and hard-stop impact.
3	Recommend without assessment, or if the highest 10% of the top-user sample (up to ten records) exceeds 40% of usage or has a concentration score over 40. Scenario ULB uses 130% of included credits per user; canonical is 5,000.	Business and budget owners confirm the enterprise default and users for whom a hard stop is acceptable.
4	Candidate users consume more than 1 times and no more than 2 times average consumption.	Cost-center owner confirms legitimate elevated work and the 6,000 ceiling.
5	Candidate users consume more than 2 times and no more than 3 times average consumption.	Cost-center owner confirms sustained high-intensity work and the 7,000 ceiling.
6	Candidate users consume more than 3 times average consumption.	Cost-center and security owners confirm exceptional need, model access, weekly review, and the 8,000 ceiling.
7	Default organization profile; also recommended when frontier-model share exceeds 30% for standard organizations.	Business, security, and license owners approve Business licensing and the example restricted feature set.
8	Organization consumption is more than 1.1 times and no more than 1.75 times the assessed organization average.	Owners approve Enterprise licensing, limited Agent Mode, repository refactor, and model access.
9	Organization consumption is more than 1.75 times and no more than 2.5 times the assessed organization average.	Owners approve architecture use cases, integrations, limited Cloud Agent, and model access.
10	Organization consumption is more than 2.5 times the assessed organization average.	Executive, security, license, and budget owners approve all-access need and enhanced monitoring.

Do not use an average-based rule when any organization, license, user, or usage source is unavailable. Resolve the gap or mark the recommendation as non-actionable. Before assignment, review business purpose, data sensitivity, security controls, license eligibility, expected cost, interruption impact, and owner acceptance.

Profile register and naming

Maintain one controlled profile register. Use stable record identifiers and version every approved change.

Required field	Purpose
Register record ID	Stable customer identifier independent of a GitHub display name
Canonical class ID and slug	Links the customer profile to Class 1 through 10
Customer profile name and version	Identifies the approved production variant
Profile type and scope	Distinguishes pool, budget, ULB, and organization policy records
Enterprise and organization	Records the assignment boundary
Cost center and team	Records the ownership and population mapping
User or cohort rule	Defines who is assigned and prevents duplicate assignment
SKU and license basis	Records license types, counts, source, and snapshot time
Credit or currency limit	Records value, unit, frequency, rate basis, and effective date

Required field	Purpose
Alert and hard-stop settings	Records 75%, 90%, and 100% behavior and recipients
Models and features	Records approved values, restrictions, and Not specified fields
Business owner	Accepts purpose and user impact
Budget and cost-center owners	Accept financial exposure and chargeback
Security and license approvers	Accept policy, integration, entitlement, and data-handling posture
Technical operator	Performs the approved production action
Evidence links	Links assessment, scenario, report, approval, change, verification, and rollback records
Review and expiry dates	Forces periodic reassessment and time-bounds exceptions
Status	Uses proposed, pilot, approved, active, suspended, retired, or rolled-back

Apply the customer rules in the [GitHub Governance Naming Standards](#). The built-in identifiers such as `aic-0011-ovr`, `org-0001-dig`, and `aic-coverage-usage-team` are examples and may not comply with customer naming patterns. Do not rename existing provider resources without an approved migration. Record noncompliance as drift or an approved exception.

Gated implementation procedure

Phase 1: Prepare

Accountable roles: service owner, enterprise budget owner, security owner, license owner, and GitHub enterprise administrator.

1. Name every accountable role and escalation delegate.
2. Define enterprise, organization, user, SKU, and billing-cycle scope.
3. Confirm application access and read permissions with non-destructive checks.
4. Approve evidence storage, data handling, retention, and secret handling.
5. Publish customer naming standards and the profile-register template.
6. Define pilot success, alert response, hard-stop communication, and rollback criteria.

Outputs and evidence:

- Signed scope and role matrix
- Access test with no credentials captured
- Approved data-handling and evidence plan
- Profile-register template
- Pilot and rollback criteria

Exit gate: all owners accept the scope, permissions, evidence location, and decision rights. No assessment begins with an unowned data source.

Phase 2: Assess and baseline

Accountable roles: service owner, license owner, budget owner, and GitHub enterprise administrator.

1. Enter validated enterprise and organization slugs and select the assessment period.
2. Run the assessment and retain its timestamp, scope, rate basis, and warnings.
3. Record GitHub-sourced usage, inventory, budget, cost-center, organization, team, user, and membership data.

4. Mark each missing source as unavailable. Do not substitute zero or an empty list.
5. Reconcile license totals, included pool, gross usage, included usage, metered usage, and budget consumption to authoritative records.
6. Label the daily trend, concentration score, governance gaps, and readiness score as synthesized or heuristic.
7. Establish the approved baseline and document unresolved variances.

Outputs and evidence:

- Dated assessment result and warning list
- Source inventory and data-quality log
- Billing and license reconciliation
- Approved baseline with variance explanations

Exit gate: required sources are complete or explicitly accepted as limitations, and budget and license owners approve the reconciled baseline.

Phase 3: Design

Accountable roles: service owner, budget owner, security owner, license owner, organization business owners, and cost-center owners.

1. Create Class 1 and Class 2 register records using the reconciled pool and customer-selected enterprise spending limit.
2. Model the Class 3 Universal ULB and Class 4 through 6 population allocations.
3. Confirm that allocations are nonnegative and total the licensed population exactly once.
4. Review advisory thresholds, then confirm each user cohort using business purpose and interruption impact.
5. Map each organization to one proposed Class 7 through 10 policy profile.
6. Review every proposed model and feature value with security and license owners.
7. Apply customer naming rules and record drift from built-in examples.
8. Compare scenarios under promotional and standard rate bases when the assessment falls within the promotional period.

Outputs and evidence:

- Versioned proposed profile register
- Scenario comparison and population mapping
- Organization policy mapping
- Security, licensing, and naming review notes

Exit gate: every proposed assignment has one owner, a documented rationale, complete inputs, and no unresolved duplicate or missing population mapping.

Phase 4: Approve

Accountable roles: enterprise budget owner, security owner, license owner, business owners, and executive sponsor for accepted exceptions.

1. Review the report and trace every recommendation to its source, calculation, or heuristic.
2. Approve the enterprise limit, ULB values, alert recipients, hard-stop response, and policy mappings separately.
3. Record approvals from business, security, license, budget, and cost-center owners as applicable.
4. Time-bound exceptions and define compensating controls, expiry, and reapproval conditions.
5. Approve pilot population, duration, support plan, success measures, and rollback trigger.

6. Open customer change records for manual production configuration.

Outputs and evidence:

- Signed profile-register version
- Decision log and exception register
- Approved pilot and rollback plan
- Authorized production change records

Exit gate: each profile is approved, rejected, or returned to design with a named action owner.

Phase 5: Pilot

Accountable roles: pilot lead, GitHub enterprise administrator, service owner, support lead, and affected business owners.

1. Capture the pre-change configuration and a tested restoration procedure.
2. Have an authorized administrator manually configure approved pilot controls in GitHub.
3. Independently verify values, scope, recipients, and hard-stop behavior against the approved register.
4. Notify pilot users of limits, alert handling, support, and interruption paths.
5. Monitor consumption, alerts, workflow impact, support demand, and exception requests each week.
6. Reconcile pilot results to billing records and collect structured user and owner feedback.
7. Roll back when a trigger is met, evidence is incomplete, or an unapproved population is affected.

Outputs and evidence:

- Before and after configuration evidence
- Independent verification record
- Weekly pilot measures and feedback
- Incident, exception, and rollback records

Exit gate: success measures are met for the approved duration, reconciliations are within accepted tolerance, and owners approve rollout. Otherwise, roll back or return to design.

Phase 6: Roll out

Accountable roles: service owner, GitHub enterprise administrator, change manager, support lead, and organization business owners.

1. Sequence rollout by organization or cohort to limit simultaneous impact.
2. Confirm the approved register version and current source data before each wave.
3. Communicate effective dates, limits, alerts, support, and exception routes.
4. Have authorized administrators manually apply the approved GitHub changes.
5. Verify each wave against the profile register and record discrepancies.
6. Pause or roll back a wave when hard-stop impact, data variance, or assignment error exceeds approved criteria.
7. Close each change only after evidence and ownership records are complete.

Outputs and evidence:

- Wave plan and communications
- Change and verification records for each wave
- Updated active profile register
- Rollout reconciliation and acceptance

Exit gate: all in-scope assignments match the active register, unresolved variances have owners and due dates, and operations accepts support ownership.

Phase 7: Operate and improve

Accountable roles: service owner, budget owner, security owner, license owner, cost-center owners, and organization business owners.

1. Run weekly operational reviews for alerts, high-consumption cohorts, hard stops, incidents, and active exceptions.
2. Run monthly financial reviews for license counts, pool estimates, consumption, overage, cost-center allocation, forecast, and reassignment.
3. Run quarterly governance reviews for profile values, organization policies, models, features, ownership, exceptions, training, and naming drift.
4. Rerun assessment and scenarios after material license, rate, organizational, policy, or usage changes.
5. Process changes through versioning, approvals, testing, implementation, verification, and rollback readiness.
6. Retire stale profiles and preserve evidence according to customer retention policy.

Outputs and evidence:

- Weekly, monthly, and quarterly review records
- Reconciliations, trends, incidents, and exception decisions
- Versioned profile changes and retirement records
- Improvement backlog with owners and target dates

Exit gate: recurring controls operated on schedule, exceptions are current, and the active register matches verified production state.

Monitoring and response controls

Apply these responses to profiles that support alerts and hard stops. Class 1 uses thresholds for pool monitoring, but `stopUsage` is `not-applicable`. Classes 2 through 6 use `stopUsage: true` in the canonical catalog.

Thre shold	Required response	Evidence
75%	Notify the profile and budget owners; validate source completeness; confirm burn rate, forecast, assignment, and remaining need.	Alert record, owner acknowledgment, refreshed forecast, and data-quality check
90%	Escalate to the cost-center or enterprise budget owner; review headroom, continuity impact, exceptions, and communications; freeze nonessential increases.	Escalation decision, impact assessment, approved action, and user communication
100%	Enforce the configured hard stop where supported; open an incident or financial-control record; notify owners and users; require emergency approval for restoration or increase.	Provider event, incident record, approval or rollback decision, and post-event reconciliation

Cadenc e	Minimum review
Weekly	75%, 90%, and 100% events; Classes 4 through 6; top-user concentration; incidents; data warnings; pilot measures; expiring exceptions
Monthly	License inventory; included pool; gross, included, and metered usage; overage; budget consumption; cost-center mapping; forecast; assignments
Quarterl y	All ten profiles; organization policy mappings; models and features; owner access; naming drift; exception necessity; training; changes

Data quality and reconciliation

Use this resolution sequence whenever a source is incomplete or totals differ:

1. Preserve the original result and warning. Do not overwrite evidence.
2. Classify the value as GitHub-sourced, calculated, heuristic, synthesized, or manual.
3. Confirm scope, timestamp, billing month, organization list, SKU mapping, and rate basis.
4. Compare enterprise totals with organization totals and identify expected coverage differences.
5. Compare license-owner records with organization and enterprise-direct license inventory.
6. Compare gross usage, included usage, net metered usage, pool use, and budget consumption without treating them as interchangeable.
7. Recalculate scenarios with reconciled inputs and the applicable rate basis.
8. Record variance amount, cause, owner, disposition, and approval.
9. Block production assignment when a missing value affects averages, hard-stop impact, license eligibility, or spending exposure.

An empty result is zero only when the source completed successfully and the provider explicitly returned no matching records. A failed, hidden, unauthorized, unsupported, or partially retrieved source is unknown or unavailable.

Exceptions, change control, and rollback

Every exception must identify the profile, affected population, business reason, risk, compensating controls, budget impact, approvers, effective date, expiry date, and rollback trigger. Review exceptions weekly and reapprove or close them before expiry.

Use this change sequence for limits, assignments, models, features, names, and alert recipients:

1. Open a change request linked to the active profile-register version.
2. Capture the current production state and reconciliation baseline.
3. Assess financial, security, license, business, support, and hard-stop impact.
4. Obtain required approvals and update the proposed register version.
5. Test through a pilot or documented impact-validation path.
6. Have an authorized administrator apply the change manually.
7. Independently verify provider state and monitor the agreed observation period.
8. Roll back on assignment error, unauthorized access, unexpected interruption, unreconciled variance, or failed success criteria.
9. Close the change after evidence is complete and the active register is updated.

Rollback restores the captured prior provider state and prior active register version. It does not delete assessment history, approval records, incidents, or evidence needed for audit.

Report and evidence checklist

Before approving a profile or closing a rollout wave, confirm the evidence set contains:

- Assessment scope, timestamp, source warnings, and authorization context
- GitHub-sourced usage, license, budget, cost-center, organization, team, user, and membership records used in the decision
- Simulator inputs, population allocation, scenario name, calculation outputs, and rate basis
- Explicit labels for application calculations, heuristics, and synthesized daily trend

- Reconciliation to authoritative license and billing records
- Profile-register version with mappings, values, owners, and effective dates
- Business, security, license, budget, and cost-center approvals as applicable
- Exception records with expiry and compensating controls
- Before and after provider configuration evidence captured by authorized operators
- Independent verification, alert tests, pilot results, and user communications
- Change, incident, rollback, and post-change monitoring records
- Generated PDF reviewed against the current profile register and standard rate basis
- Post-August 31, 2026 review for decisions that used promotional assessment rates

Glossary

Term	Definition
AI Credit (AIC)	Unit used by the application to represent GitHub Copilot usage-based consumption.
Included pool	Estimated credits included with configured license inventory for the applicable rate basis.
Metered overage	Consumption beyond the included pool that the application estimates at 0.01 USD per credit.
User-Level Budget (ULB)	Monthly per-user credit ceiling used by Universal and cost-center profiles.
Universal ULB	Enterprise-wide default ULB, represented by canonical Class 3 at 5,000 credits per user per month.
Cost-center ULB override	More specific Class 4, 5, or 6 per-user ceiling for an approved mapped cohort.
Enterprise Spending Limit	Tier 1 hard cap on metered overage across the enterprise.
Organization policy profile	Class 7 through 10 planning record for SKU, model, and feature settings. It is not an organization budget.
Hard stop	Configured behavior that prevents further applicable usage when a supported limit reaches 100%.
Advisory heuristic	Application rule that identifies a candidate action but requires human validation and approval.
Synthesized data	Display or planning data generated from other values rather than retrieved as raw provider telemetry.
Profile register	Controlled record of profile values, mappings, ownership, approvals, evidence, status, and version history.
Reconciliation	Comparison that explains differences among source records, calculations, scenarios, and billing evidence.
Rate basis	Credit-per-license values used for a pool estimate, including time-limited promotional or standard rates.

References

- [GitHub Governance Naming Standards](#)
- The application Assessment page for GitHub-sourced results and data warnings
- The application Simulator page for manually confirmed planning assumptions and scenarios
- The application Recommendations page for advisory heuristic output
- The application Report page for preview and PDF generation