

COLOR E-PAPER EDITION

Enterprise AI Credits Governance Implementation Guide

Design, approve, pilot, roll out, and operate configurable AI Credits governance controls.

Contents

Audience, scope, and outcomes

Roles and prerequisites

1. Get AI Included Credits

2. Set Enterprise Spending Budget

3. Establish User Level Budget (ULB) Profiles

4. Create a taxonomy for ULBs

4.1 Example ULB Budget Profiles Taxonomy

5. Create Enterprise Team Naming Standards

5.1 Enterprise Teams Naming Standard Examples

6. Create or Update Enterprise Team Names or Properties

7. Establish Cost Center Naming Standards

7.1 Cost Center Naming Standard Examples

8. Create or Update ULB-based Cost Center Names or Properties

9. Establish Organization based Cost Center Naming Standards

9.1 Organization based Cost Center Naming Standard Examples

10. Create or Update Organization based Cost Center Names or Properties

11. Assess Current GitHub Copilot Policies

12. Define Desired State of GitHub Copilot Policy Set Profiles

13. Define Budgets

14. Monitor Usage

15. Analyze Monthly AI Credits Usage

16. Obtain Qualitative User Feedback

17. Review and Adjust Policy Sets

18. Review and Adjust Budgets

19. Token Optimization Training

20. Continuous Improvement and Iteration (Sections 14-19)

Appendix A: Design and assignment criteria

Appendix B: Governance architecture and workflow

Appendix C: Lifecycle phases

Appendix D: Monitoring and response

Appendix E: Reconciliation and data quality

Appendix F: Exceptions, change control, and rollback

Appendix G: Review cadence

Appendix H: Getting-started checklist

Appendix I: Evidence checklist

Appendix J: Glossary

Use current authoritative provider records and controlled tests before relying on rates, entitlements, alerts, precedence, or enforcement behavior.

Audience, scope, and outcomes

This guide is for customer budget owners, service owners, security teams, license managers, business owners, administrators, and assurance reviewers who govern usage-based AI services.

No software integration, source-code package, or particular assessment tool is required. Customers may use approved provider portals, exports, finance records, spreadsheets, scripts, or other approved tools.

AI Credit (AIC) means the provider-defined unit used to measure eligible AI consumption. User-Level Budget (ULB) means a customer-defined per-user spending or consumption boundary where the provider supports one.

The governance process should produce:

- A defined scope and accountable ownership model
- A reconciled consumption, entitlement, and cost baseline
- Configurable financial, user, cohort, and access-policy controls
- Approved profile assignments with evidence and review dates
- Tested alert, enforcement, exception, and rollback procedures
- An operating cadence that keeps records aligned with verified provider state

[!IMPORTANT] Confirm current rates, entitlements, data availability, supported controls, control scope, precedence, alerts, and enforcement behavior from authoritative provider records and documentation. Use controlled tests to verify behavior environment before relying on any control.

Roles and prerequisites

Role	Accountability
Executive sponsor	Accepts risk posture and material unresolved exceptions
Enterprise budget owner	Approves financial exposure, response options, and variance
Service owner	Owns the method, registers, evidence, reviews, and control health
Security owner	Approves access posture, data handling, and compensating controls
License owner	Validates eligibility, entitlement, inventory, and assignment
Business owner	Confirms need, affected population, duration, and interruption risk
Cohort or cost owner	Accepts cohort membership, funding, and review obligations
Authorized administrator	Applies and verifies approved provider changes
Assurance reviewer	Reviews evidence, approvals, reconciliation, and control operation

Complete these prerequisites before design begins:

- Name role owners and escalation delegates
- Define enterprise, organizations, AI Credit SKUs, population, and period scope
- Approve evidence storage, retention, access, and data handling
- Confirm access to authoritative provider, license, and finance records
- Establish change, incident, exception, and rollback processes
- Define acceptable reconciliation tolerance and unresolved-data treatment
- Document planned controlled tests

1. Get AI Included Credits

Retrieve the current AI included credits for the enterprise.

AI included credits are the provider-defined units bundled with eligible licenses or services for a specific billing period. Establish the enterprise total from current license inventory, applicable SKU rates, assignment dates, and authoritative provider records. Record both the total and the assumptions used to calculate or retrieve it. *Tip: Navigate to: Enterprise->Billing and licensing (tab)->AI Usage Pane*

This value is the baseline for forecasting when metered overage could begin. It also allows budget owners to distinguish consumption already covered by an entitlement from consumption that may create an additional charge.

Included credits are not a cash budget, and a displayed balance may not represent every organization, SKU, adjustment, or delayed usage record. Rates, pooling rules, expiration behavior, and billing-period boundaries can change. Treat unavailable or partial values as unknown and reconcile the total before using it for a financial decision.

2. Set Enterprise Spending Budget

Set the enterprise spending budget for AI credits as a universal buffer.

The enterprise spending budget is the approved aggregate boundary for metered AI Credit charges after included entitlements are exhausted. Define its scope, currency or credit unit, billing period, alert thresholds, recipients, decision owners, and the response expected when each threshold is reached.

This control provides a financial backstop across the enterprise. It limits unplanned exposure, gives owners time to investigate abnormal consumption, and establishes clear authority for increasing, pausing, or retaining the limit.

A provider budget may be informational, alert-only, or enforcing. Do not assume that reaching the displayed amount stops usage. Confirm whether prior commitments, delayed events, excluded products, or organization-specific spend sit outside the boundary, and test enforcement and rollback behavior before relying on it.

Edit Copilot AI Credits

SKU

All AI Credit SKUs

A budget for all SKUs that consume AI credits, where [overages spend is enabled](#).

Supported SKUs:

- Copilot AI credits
- Cloud Agent AI credits
- Spark AI credits

Budget scope

Enterprise

Scoped budgets track spending for all organizations, repositories and cost centers in the enterprise

 autocloudarc-labs

☐ Exclude cost center usage from this budget [Preview](#)

If selected, cost center usage will not be limited by this budget. [Learn more](#).

Budget

Set a budget amount to track your spending on a monthly basis

Budget amount

\$ 60

☒ Stop usage when budget limit is reached

Spending won't exceed your set budget

Alerts

Receive emails when usage reaches 75%, 90%, and 100% of the budget amount

☒ Receive budget threshold alerts

Alert recipients

Account owners and billing managers receive budget alert emails by default. You can add other recipients if needed.

Select recipients ▾

1 selected

 autocloudarc

×

Figure: 1 - Enterprise spending budget configuration.

3. Establish User Level Budget (ULB) Profiles

Define user level budget profiles for AI credits.

A User-Level Budget profile is a reusable per-user consumption boundary with a defined period, value, alert behavior, enforcement posture, ownership, and eligible population. Establish an enterprise-wide default and, where justified, more specific profiles for approved cohorts with materially different needs.

Profiles make per-user treatment consistent and reviewable. They reduce the risk that one account silently consumes a disproportionate share of the pool while providing a documented path for legitimate higher-usage work.

A ULB is a financial control, not proof of license eligibility, feature access, or business need. More specific assignments may override the universal default, and overlapping assignments can produce unexpected precedence. Confirm how the provider evaluates limits, alerts, resets, and hard stops, then document exceptions and interruption impact before assignment.

Edit Copilot AI Credits

SKU

All AI Credit SKUs

A budget for all SKUs that consume AI credits, where [overages spend is enabled](#).

Supported SKUs:

- Copilot AI credits
- Cloud Agent AI credits
- Spark AI credits

Budget scope

Users

Each user in your enterprise gets their own individual budget

 autcloudarc-labs

Budget

Set a budget amount to track your spending on a monthly basis

Budget amount

\$ 50 per user

5000 AI credits (\$0.01/credit)

- ☒ Stop usage when user's budget limit is reached ⓘ
User spending will not exceed your set per user budget.

Update budget

Cancel

Figure: 2 - User-universal budget configuration.

Edit ccr-0011-ovr

SKU

All AI Credit SKUs

A budget for all SKUs that consume AI credits, where [overages spend is enabled](#).

Supported SKUs:

- Copilot AI credits
- Cloud Agent AI credits
- Spark AI credits

Budget scope

Cost center (Per-user)

Applies to every licensed user in a selected cost center. Overrides the Universal budget for those users.

 ccr-0011-ovr

Budget

Set a budget amount to track your spending on a monthly basis

Budget amount

per user

6000 AI credits (\$0.01/credit)

☒ Stop usage when user's budget limit is reached 

User spending will not exceed your set per user budget.

Update budget

Cancel

Figure: 3 - Cost-center per-user budget configuration for ccr-0011-ovr.

13 / 72

Edit ccr-0012-abd

SKU

All AI Credit SKUs

A budget for all SKUs that consume AI credits, where [overages spend is enabled](#).

Supported SKUs:

- Copilot AI credits
- Cloud Agent AI credits
- Spark AI credits

Budget scope

Cost center (Per-user)

Applies to every licensed user in a selected cost center. Overrides the Universal budget for those users.

 ccr-0012-abd

Budget

Set a budget amount to track your spending on a monthly basis

Budget amount

\$ 70 per user

7000 AI credits (\$0.01/credit)

- ☒ Stop usage when user's budget limit is reached ⓘ
User spending will not exceed your set per user budget.

Update budget

Cancel

Figure: 4 - Cost-center per-user budget configuration for ccr-0012-abd.

Edit ccr-0013-exp

SKU

All AI Credit SKUs

A budget for all SKUs that consume AI credits, where [overages spend is enabled](#).

Supported SKUs:

- Copilot AI credits
- Cloud Agent AI credits
- Spark AI credits

Budget scope

Cost center (Per-user)

Applies to every licensed user in a selected cost center. Overrides the Universal budget for those users.

 ccr-0013-exp

Budget

Set a budget amount to track your spending on a monthly basis

Budget amount

\$ 80 per user

8000 AI credits (\$0.01/credit)

☒ Stop usage when user's budget limit is reached ⓘ
User spending will not exceed your set per user budget.

Update budget

Cancel

Figure: 5 - Cost-center per-user budget configuration for ccr-0013-exp.

4. Create a taxonomy for ULBs

Define a taxonomy for user level budget profiles.

A ULB taxonomy is the controlled classification scheme used to name and distinguish profiles. Define categories by business need, expected consumption range, interruption tolerance, duration, funding owner, and review cadence rather than by usage volume alone.

The taxonomy gives reviewers a shared vocabulary for comparing profiles and helps administrators apply the same criteria across organizations. It also supports reporting, exception management, and lifecycle decisions without embedding arbitrary amounts or temporary organizational details in every name.

Taxonomy labels do not automatically classify users or enforce a limit. Avoid terms that imply performance, seniority, or permanent entitlement, and do not promote users solely because they crossed a threshold once. Define entry, exit, expiry, and reassessment criteria so categories do not become stale privileges.

4.1 Example ULB Budget Profiles Taxonomy

user-universal-50

The default per-user budget for all eligible users, with a monthly limit of 50 AI Credits.

user-cost-center-ccr-0011-ovr-60 A cost-center-specific profile for the "ccr-0011-ovr" cohort where "ccr" is the cost center prefix and "ovr" represents "overage usage" limit of 6,000 AI Credits.

user-cost-center-ccr-0012-abd-70 A cost-center-specific profile for the "ccr-0012-abd" cohort where "ccr" is the cost center prefix and "abd" represents "abundant usage" limit of 7,000 AI Credits.

user-cost-center-ccr-0013-exp-80 A cost-center-specific profile for the "ccr-0013-exp" cohort where "ccr" is the cost center prefix and "exp" represents "exponential usage" limit of 8,000 AI Credits.

5. Create Enterprise Team Naming Standards

Define enterprise team naming standards in order to ensure consistency and clarity across the organization.

An enterprise team naming standard is a documented grammar for team names and stable identifiers. Define required segments such as service, business unit, purpose, environment, or governance class, together with approved values, character rules, length limits, examples, and an accountable owner.

Consistent names let administrators and reviewers understand a team's intended scope without relying on personal knowledge. They improve search, automation, reporting, assignment reviews, and incident response, especially when similar teams exist across many organizations.

A name is metadata, not evidence of membership, ownership, entitlement, or policy enforcement. Provider restrictions and existing integrations may limit which fields can change. Prefer stable identifiers for automation, preserve aliases or migration mappings when names change, and avoid embedding sensitive information or values likely to become obsolete.

5.1 Enterprise Teams Naming Standard Examples




Figure: 6 - Enterprise team naming standard examples with cost-center and organization mappings.

Add a clear description to every team. A description is a human-readable statement of the team's purpose, scope, and owner. It helps reviewers understand the intended membership and use of the team without relying on personal knowledge. It also improves search, reporting, and assignment reviews. For billing, consider including the cost center or organization in the description to make it easier to reconcile with finance records.

Edit Enterprise team

Name *

ccr-abundant-usage-team

23/100

Description

Greater business investment of AI credits for super users to drive more impact.

associated-cost-center: aic-0012-abd

organization: all

136/350

What is this team all about? The maximum length is 350 characters

Team access



All organizations

10



Notification setting

Choose whether team members receive notifications when this team is @mentioned.



Enabled

Everyone will be notified when the team is @mentioned.



Disabled

No one will receive notifications.

Update team

Cancel

Additional options

Figure: 7 - Enterprise team configuration for ccr-abundant-usage-team.

6. Create or Update Enterprise Team Names or Properties

Define or update enterprise team names or properties according to the standards established above.

Create new enterprise teams or bring existing team names, descriptions, ownership fields, and other approved properties into alignment with the naming standard. Record the provider identifier, prior value, proposed value, owner, reason, affected assignments, and effective date for every change.

Applying the standard turns a design convention into an operable inventory. It also creates reliable references for ULB cohorts, cost centers, policy profiles, communications, and assurance evidence.

Renames and property changes can break scripts, reports, synchronization, or manual procedures that use a mutable name as a key. Membership changes may also alter access or financial treatment. Assess those effects separately, stage high-impact changes, verify the resulting membership and assignments, and retain a tested restoration path.

7. Establish Cost Center Naming Standards

Define cost center naming standards in order to ensure consistency and clarity across the enterprise.

A cost center naming standard defines how billing and governance groupings are identified across the enterprise. Specify which segments represent the service, funding owner, business unit, environment, purpose, or profile class, and define how names map to authoritative finance identifiers.

The standard makes allocation and reconciliation repeatable. A reviewer should be able to connect a provider cost center to its owner, approved population, budget treatment, and finance record without interpreting an informal label.

The provider name should not replace the finance system's immutable identifier, and creating a label does not establish charge-back authority or complete data coverage. Account for provider length and character constraints, mergers and renames, shared funding, and one-to-many mappings between technical and finance structures.

7.1 Cost Center Naming Standard Examples

ccr-0010-uni

Universal cost center for general-purpose AI Credit usage, where "ccr" is the cost center prefix and "uni" represents "universal usage".

ccr-0011-ovr

Cost center for the "ccr-0011-ovr" cohort where "ccr" is the cost center prefix and "ovr" represents "overage usage".

ccr-0012-abd

Cost center for the "ccr-0012-abd" cohort where "ccr" is the cost center prefix and "abd" represents "abundant usage".

ccr-0013-exp

Cost center for the "ccr-0013-exp" cohort where "ccr" is the cost center prefix and "exp" represents "exponential usage".

8. Create or Update ULB-based Cost Center Names or Properties

Define or update ULB-based cost center names or properties according to the standards established above.

A ULB-based cost center represents an approved cohort whose members receive a specific per-user budget profile. Create or update its name, owner, mapped team or population, profile reference, funding context, expiry, and review date in accordance with the cost center standard.

This mapping makes higher or otherwise specialized per-user limits traceable. It separates documented exceptions from the universal ULB and gives budget owners a defined population to review, forecast, and reconcile.

This is still a collection of per-user controls, not necessarily a pooled cost center budget. Membership errors or overlapping mappings can assign the wrong limit, and a provider may use different precedence rules than the design assumes. Verify effective assignments after every change and remove users when their approved need or time-bound exception ends.

9. Establish Organization based Cost Center Naming Standards

Define organization based cost center naming standards in order to ensure consistency and clarity across the enterprise.

An organization-based cost center naming standard identifies billing groupings whose primary scope is a provider organization, license SKU, or accountable business owner. Define how the name distinguishes organization allocation from ULB cohort allocation and how it links to the authoritative organization and finance identifiers.

The distinction supports SKU-level forecasting, ownership, and reconciliation without confusing an organizational boundary with a user-level exception. It also helps reviewers identify whether a variance arises from license inventory, organization membership, usage, or an allocation rule.

Organizations and finance cost centers rarely have a permanent one-to-one relationship. Shared services, user transfers, multiple SKUs, and reorganizations can create one-to-many or many-to-one mappings. Define effective dates and mapping ownership, and do not infer a provider budget or policy merely from the cost center name.

9.1 Organization based Cost Center Naming Standard Examples

org-0001-dig

Organization-based cost center for the "org-0001-dig" organization where "org" is the cost center prefix and "dig" represents "digital services".

org-0002-fin

Organization-based cost center for the "org-0002-fin" organization where "org" is the cost center prefix and "fin" represents "financial services".

org-0003-hlt

Organization-based cost center for the "org-0003-hlt" organization where "org" is the cost center prefix and "hlt" represents "healthcare solutions".

org-0004-mfg

Organization-based cost center for the "org-0004-mfg" organization where "org" is the cost center prefix and "mfg" represents "manufacturing group".

org-0005-mar

Organization-based cost center for the "org-0005-mar" organization where "org" is the cost center prefix and "mar" represents "marine industries".

org-0006-mil

Organization-based cost center for the "org-0006-mil" organization where "org" is the cost center prefix and "mil" represents "military operations".

org-0007-opn

Organization-based cost center for the "org-0007-opn" organization where "org" is the cost center prefix and "opn" represents "open source initiatives".

org-0008-rtl

Organization-based cost center for the "org-0008-rtl" organization where "org" is the cost center prefix and "rtl" represents "retail operations".

org-0009-spc

Organization-based cost center for the "org-0009-spc" organization where "org" is the cost center prefix and "spc" represents "space fleet".

org-0010-prp

Organization-based cost center for the "org-0010-prp" organization where "org" is the cost center prefix and "prp" represents "property management".

10. Create or Update Organization based Cost Center Names or Properties

Define or update organization based cost center names or properties according to the standards established above.

Create or update organization-based cost centers with their approved names, organization references, SKU scope, finance mapping, owner, and review date. Capture the prior state and reconcile organization membership and license inventory before applying the change.

Maintaining these objects provides an auditable allocation layer between provider billing records and internal financial accountability. It allows enterprise totals to be explained by organization and SKU while preserving a named owner for unresolved variances.

Allocation changes may apply only to future usage and may not restate historical charges. Moving an organization or user can also change reports without changing actual consumption. Confirm provider effective-date behavior, avoid double counting users or usage across mappings, and verify both provider and finance records after the update.

11. Assess Current GitHub Copilot Policies

Retrieve the current GitHub Copilot policies for the enterprise and organizational levels.

Assessing current GitHub Copilot policies means creating a timestamped baseline of configured and effective settings at every relevant scope. Capture the source, retrieval method, permissions, enterprise defaults, organization values, exceptions, unsupported fields, and any setting whose effective state cannot be confirmed.

The baseline is necessary for gap analysis, impact assessment, rollback, and independent verification. Without it, a desired-state change can unintentionally remove an existing restriction, preserve hidden drift, or misidentify an inherited value as an organization decision.

An unavailable value is not equivalent to disabled or unrestricted. Portals, APIs, exports, and documentation may expose different portions of the policy surface, and the effective result may depend on SKU, inheritance, precedence, or propagation delay. Reconcile sources and preserve unknowns explicitly.

11.1 Assess Enterprise GitHub Copilot Policy Set Profile

Retrieve the current enterprise GitHub Copilot policy set profile.

The enterprise policy assessment identifies the highest-scope Copilot defaults, constraints, assignments, and permitted organization variation. Record each setting's configured value, source, owner, affected population, and observed effective behavior where a controlled test is appropriate.

This assessment defines the control envelope within which organization profiles operate. It reveals broad-impact settings, establishes the rollback baseline, and prevents organization-level designs from depending on behavior the enterprise layer prohibits.

The provider may not expose a single object called an enterprise policy set profile. The profile can therefore be a customer-controlled record that maps to several provider settings. Do not fill unsupported or inaccessible fields with assumed defaults, and verify whether changes affect existing organizations, future organizations, or both.

11.2 Assess Organizational GitHub Copilot Policy Set Profile

Retrieve the current organizational GitHub Copilot policy set profile.

The organization policy assessment records each in-scope organization's local Copilot settings, inherited enterprise values, approved exceptions, eligible SKUs, and effective user population. Compare organizations using the same field definitions so omissions and drift are visible.

Organization-level evidence is important because a compliant enterprise default does not guarantee a consistent effective posture. Local overrides, historical configuration, licensing differences, or incomplete rollout can produce materially different access and risk outcomes.

A displayed organization value may be inherited, locked, stale, or inapplicable to some users. Do not treat configuration equality as proof of equal behavior. Record inheritance and precedence separately, preserve inaccessible results as unknown, and validate representative organizations before generalizing.

11.3 Example Enterprise GitHub Copilot Policy Set Profile Schema

Use the workbook schema as an editable policy catalog and governance baseline.

The schema below is extracted from [github-copilot-policy-sets.xlsx](#), whose source metadata was verified on August 5, 2026. The workbook states that it is a reference catalog, not an API export of a specific enterprise or organization. Reconcile its values with current provider state before using it as an approval or change record.

Workbook structure

Worksheet	Purpose	Data rows	Observed scope	Observed statuses
notes-sources-Legend	Purpose, conventions, source links, conflict guidance, and budget-impact legend	Reference content	Not applicable	Not applicable

Worksheet	Purpose	Data rows	Observed scope	Observed statuses
ent-ghcp-policy-set-labs	Enterprise policy baseline for auto-cloudarc-labs	24	ent	2 enabled , 22 let-orgs-decide
org-ghcp-policy-set-digi-svcs	Organization policy baseline for autocloudarc-digital-services	28	org	15 enabled , 13 disabled

Each policy worksheet contains one row per policy and uses the same 13-column shape. The owner column changes from `ent-name` to `org-name` according to scope. The first ten fields are populated in every current policy row. The three change-tracking fields are present but empty in the supplied baseline.

Policy record fields

Source field	Type	Baseline population	Definition and constraints
index	Text	Required	Three-digit row identifier such as 001 ; unique within each worksheet
ent-name or org-name	Text	Required	Enterprise or organization that owns the policy record; field must agree with scope
policy-set-name	Text	Required	Stable name that groups the policy rows into one profile
policy-name	Text	Required	Human-readable provider policy name; unique within each supplied profile
scope	Enum	Required	Excel list allows ent or org ; current enterprise and organization sheets use their corresponding value
description	Text	Required	Plain-language statement of the capability governed by the policy
status	Enum	Required	Excel list allows enabled , disabled , or let-orgs-decide

Source field	Type	Baseline population	Definition and constraints
settings	Text	Required	Human-readable allowed values, dependencies, surface coverage, and policy-specific caveats; not a machine-readable settings object
conflict-behavior	Text	Required	Provider precedence or cross-organization behavior; records “not specified” when the provider conflict table is silent
budget-impact	Enum	Required	Qualitative planning value: low , medium , or high ; not a price or forecast
change-date	Date	Optional	Date of an approved change; cells use the workbook display format d-mmm-yy
change-reason	Text	Optional	Business, risk, support, or compliance rationale for changing the baseline
notes	Text	Optional	Supporting evidence, exception details, limitations, or follow-up actions

Treat (scope, owner name, policy-set-name, policy-name) as the logical record key when combining worksheets. The three-digit index is a local display sequence and is not globally unique. Keep settings as explanatory evidence; store provider-specific structured values separately if automation requires them.

Controlled-value interpretation

Field	Value	Meaning
status	enabled	The policy permits or activates the capability at the recorded scope
status	disabled	The policy blocks or deactivates the capability at the recorded scope
status	let-orgs-decide	The enterprise delegates the effective choice to organizations where supported
budget-impact	high	The capability may materially increase metered consumption
budget-impact	medium	The capability normally contributes Copilot or premium-request consumption

Field	Value	Meaning
budget-impact	low	The capability is primarily governance-oriented or has negligible usage impact

Conflict behavior is policy-specific. The workbook uses least-restrictive, most-restrictive, and “not specified” descriptions rather than reducing all policies to one precedence rule. For users licensed through multiple enterprises, the notes state that the most restrictive enterprise policy usually applies, with exceptions for AI Credit paid usage and GitHub Spark.

Source references

Subject	Provider reference
Policy concepts	About policies for GitHub Copilot
Enterprise policy management	Manage policies and features for Copilot in your enterprise
Organization policy management	Manage policies for Copilot in your organization
Policy conflicts	Policies for Copilot in multiple organizations and enterprises
Supported policy surfaces	Copilot policy support by feature and client

[!CAUTION] The enterprise worksheet is named `ent-ghcp-policy-set-labs`, but every current `policy-set-name` value is `org-ghcp-policy-set-autocloudarc-labs`. Confirm whether that `org-` prefix is intentional before using it as a stable enterprise profile identifier.

The workbook's `low,medium,high` list validation extends from `budget-impact` into `change-date` cells. On the organization sheet it covers only change-date rows 2 through 25. Correct those validation ranges before entering dates; otherwise Excel may reject valid dates or apply inconsistent validation across policy rows.

12. Define Desired State of GitHub Copilot Policy Set Profiles

Establish the desired state of GitHub Copilot policy set profiles for the enterprise and organizational levels.

The desired state is the approved target for Copilot feature, model, service, integration, and access posture at enterprise and organization scopes. Define each target value with its business rationale, risk owner, eligible population, source requirement, exception path, validation method, and review date.

A documented target converts broad governance intent into testable decisions. It supports consistent implementation, exposes the gap from current state, and lets approvers evaluate user impact before an administrator changes provider configuration.

Design only for capabilities confirmed in current provider documentation and the applicable SKUs. Keep access-policy profiles separate from financial budgets and ULBs even when they share populations. Where the provider has no named profile object, maintain the profile in the controlled register and map it to the individual settings that implement it.

Note: You may reuse the same schema above for enterprise and organizational desired-state profiles.

12.1 Enterprise Profile

Set the desired state of the enterprise GitHub Copilot policy set profile.

The enterprise profile is the default or mandatory access posture intended to apply across the enterprise. It should define which decisions are centralized, which may be delegated to organizations, and what minimum controls no local profile may weaken.

This profile creates a stable baseline for security, compliance, support, and user experience. It reduces accidental variation and makes organization exceptions explicit rather than allowing historical settings to become policy.

Enterprise settings have a wide blast radius and can interrupt critical work. Do not select restrictive or permissive values solely for consistency. Assess license compatibility, data handling, operational dependencies, and fallback options, then validate representative use cases before broad application.

12.1.1 Naming Standards

Create a naming standard for Enterprise GitHub Copilot Policy set profiles that can be documented and referenced.

Define a stable naming pattern for enterprise profiles that communicates scope, posture, lifecycle state, and version without encoding volatile settings in the name. Document approved abbreviations, examples, uniqueness rules, and the relationship between a display name and the profile register identifier.

Consistent names allow approvals, tests, changes, reports, and rollback records to reference the same design unambiguously. Versioning also makes it possible to compare an approved proposal with the profile actually applied.

The name does not enforce the profile and may not exist as a provider field. Respect provider character and length restrictions, avoid confidential details, and never reuse an identifier for a materially different posture. Preserve the history of retired names and versions.

*Note: A simple approach may be to use `ent-ghcp-policy-set-
<enterprise-name>` for enterprise profiles.*

12.1.2 Policy Settings

Define the policy settings for the enterprise GitHub Copilot policy set profile.

Specify the enterprise target for each supported Copilot capability, including feature availability, model access, agent or automation modes, integrations, data-sensitive functions, and organization delegation. For every setting, record allowed values, rationale, owner, dependencies, validation evidence, and fallback behavior.

This level of detail turns the profile into an implementable and testable control rather than a descriptive label. It also helps support teams explain why a capability is available, limited, or disabled.

Settings can vary by SKU, organization, client, or rollout status, and similar labels may not have identical effects. Confirm precedence and effective behavior instead of relying on display text. Policy settings govern access posture; they must not be presented as spending limits unless the provider explicitly links the behaviors and that linkage has been verified.

12.2 Organizational Profiles

Set the desired state of the organizational GitHub Copilot policy set profile.

Organizational profiles define approved variations from the enterprise baseline for populations with different business functions, risk exposure, licenses, or operational needs. Each profile should state its qualifying criteria, owner, scope, duration, enterprise dependencies, and exception process.

These profiles allow local needs to be met without abandoning enterprise governance. A small, reusable catalog is easier to approve, support, test, and review than one bespoke configuration per organization.

An organization profile cannot override an enterprise restriction unless the provider and governance model explicitly permit it. Too many profiles increase drift and support cost, while overly broad profiles can grant unnecessary capability. Require a material rationale for every variation and retire profiles that no longer serve a distinct need.

12.2.1 Naming Standards

Create a naming standard for organizational GitHub Copilot policy set profiles that can be documented and referenced.

Define an organizational profile naming pattern that identifies the reusable posture and, where necessary, its SKU, risk class, or lifecycle state. Keep the profile name separate from the organization name so one approved design can be assigned consistently to several organizations.

The standard improves inventory, assignment review, change sequencing, and drift detection. It also prevents minor wording differences from hiding materially identical or conflicting profiles.

Names should not imply that an organization is compliant or that settings are currently applied. Maintain the actual assignment and verification state in the profile register, preserve prior versions, and account for provider constraints if the name must also be represented in an external system.

Note: A simple approach may be to use `org-ghcp-policy-set-<organization-name>` for organizational profiles.

12.2.2 Policy Settings

Define the policy settings for the organizational GitHub Copilot policy set profile.

For each organizational profile, define the local target values and identify which settings are inherited, overridden, prohibited, or not applicable. Link every deviation to an approved business need, license basis, risk decision, test, owner, and review date.

Explicit treatment of inheritance prevents reviewers from mistaking an omitted field for an intentional value. It also supports a deterministic comparison between enterprise baseline, organization proposal, and observed effective state.

Do not copy every enterprise value into every organization record if inheritance is authoritative; duplication can conceal later drift. Conversely, do not rely on inheritance without verifying it. Feature availability and effective behavior may differ by SKU or rollout, so test representative users without treating a single successful result as universal coverage.

12.3 Apply GitHub Copilot Policy Set Profiles

Apply the desired state of the GitHub Copilot policy set profiles for the enterprise and organizational levels.

Applying profiles is the controlled translation of approved register values into provider configuration. Link the change to the current baseline, confirm authority, capture prior state, sequence the affected scopes, apply only approved values, and independently verify the resulting effective behavior.

This procedure closes the gap between documented intent and operational state while preserving accountability and rollback evidence. A staged rollout limits the effect of an incorrect precedence assumption or unsupported setting.

Provider updates may be asynchronous, partially accepted, or unavailable through the chosen interface. A successful request is not proof that every user received the target behavior. Record

per-scope results, propagation windows, warnings, and exceptions, and stop the rollout when verification or success criteria fail.

12.3.1 Apply Enterprise Profile

Apply the desired state of the enterprise GitHub Copilot policy set profile.

Apply the enterprise profile through an authorized administrator after approval, impact assessment, and a representative controlled test. Change the smallest practical set of values, retain the prior configuration, and verify enterprise defaults as well as the behavior of organizations with existing local settings.

Enterprise application establishes the baseline on which all later organization work depends. Performing it first, when precedence requires that order, makes subsequent deviations visible and reduces ambiguity during verification.

Because the blast radius is broad, do not combine unrelated enterprise changes or assume rollback is instantaneous. Confirm communications, support readiness, active sessions, propagation timing, and emergency exceptions before the change, then reconcile the provider state with the approved profile register.

12.3.2 Apply Organizational Profiles

Apply the desired state of the organizational GitHub Copilot policy set profile.

Assign and apply approved organizational profiles in controlled waves. For each organization, confirm the intended profile version, enterprise dependencies, SKU, owner, affected

population, exception status, and prior state before an administrator changes configuration.

Wave-based application makes comparison and pause decisions possible. It allows the team to learn from representative organizations before exposing the entire population and produces clear evidence for each assignment.

Do not infer success from the enterprise baseline or from another organization's result. Local history, licensing, and overrides can change the outcome. Verify configured and effective state independently, record partial failures, and keep an organization on its prior approved profile until its change meets the defined exit criteria.

13. Define Budgets

Define enterprise user level and organizational budgets based on the previously established teams, cost centers, and organizational structures and configurations.

Define the approved financial control hierarchy using the established teams, cost centers, organizations, SKUs, and profile assignments. Distinguish the included-credit baseline, enterprise metered-overage boundary, universal per-user limit, cost-center ULB overrides, and organization or SKU allocations.

Layered budgets answer different questions: how much entitlement is available, how much aggregate overage the enterprise accepts, how much one user may consume, and which owner is accountable for an allocation. Documenting those layers prevents a single value from being used for incompatible decisions.

Provider support, units, scope, precedence, reset periods, alerts, and enforcement can differ by control. Do not add per-user limits together and call the result an enterprise cap, or count the same usage in several allocation views. Reconcile every proposed layer to authoritative billing behavior and keep organization policy profiles separate from financial budget records.

13.1 Enterprise Level

Set the enterprise level budgets for AI credits.

Enterprise-level budgeting combines the aggregate metered-overage boundary with the default per-user ULB. Define their common billing period, applicable products and organizations, owners, thresholds, responses, exception authority, and relationship to the included-credit pool.

Together these controls address both systemic and concentration risk. The aggregate limit constrains total exposure, while the per-user default reduces the chance that one user consumes a disproportionate amount before enterprise owners can respond.

The two controls are not interchangeable and may be evaluated at different times. A user can reach a ULB while enterprise capacity remains, or enterprise spend can reach its boundary while individual users remain below theirs. Test both paths and define which response takes precedence when several thresholds are reached together.

13.1.1 Enterprise Spend Budget

Configure the enterprise spend budget for AI credits.

Configure the enterprise spend budget as the approved limit or monitoring threshold for aggregate metered AI Credit charges. Record the amount, currency or credit basis, period, included and excluded services, alert recipients, enforcement selection, approval authority, and emergency decision path.

This budget is the final financial backstop after included credits and more specific usage controls are considered. It provides a clear boundary for forecasting, escalation, and executive risk acceptance.

The amount may not equal the total provider invoice because licenses, commitments, adjustments, or other products can sit outside its scope. Alerts may lag consumption, and a stop setting may affect supported metered activity without reversing charges already incurred. Validate the exact boundary and retain response lead time below the limit.

13.1.2 User-Universal (Universal ULB)

Configure the user-universal budget for AI credits.

Configure the Universal ULB as the default monthly per-user AI Credit boundary for every eligible user not assigned a more specific approved profile. Define the value, reset period, alerts, enforcement posture, included population, exclusions, and exception route.

The universal profile establishes equitable baseline treatment and limits concentration risk without requiring an individual decision for every user. It also makes higher-usage exceptions visible because they must be assigned to a separate profile.

A universal value should not be selected from the average alone; legitimate workloads and interruption impact can vary widely. Confirm whether the provider counts all models and features toward the same boundary, how mid-period assignments behave, and whether the ULB stops usage or only signals. Monitor for users whose approved work is repeatedly interrupted.

13.2 User-Cost-Center Level (Overage ULBs)

Create the user-cost-center level budgets for AI credits.

User-cost-center ULBs are more specific per-user profiles for approved cohorts whose needs differ materially from the universal default. Define each profile's per-user value, qualifying criteria, mapped cost center and team, funding owner, duration, alerts, enforcement, and reassessment date.

These profiles provide a governed exception path for higher-usage or specialized work while preserving accountability. They can also support a lower boundary where interruption tolerance or risk warrants tighter control.

An overage ULB is not a shared pool that any cohort member may consume, and the word "overage" does not itself authorize additional spend. Confirm that the more specific profile overrides rather than adds to the universal ULB, prevent overlapping cohort assignments, and review membership frequently enough to remove stale exceptions.

13.3 Organization (SKU based)

Develop organization (SKU based) budgets for AI credits.

Organization and SKU-based budgets allocate or monitor expected AI Credit consumption for an accountable organization and license population. Define the SKU, license count, included-credit basis, forecast, metered-overage tolerance, owner, reporting period, and reconciliation method for each allocation.

This view explains enterprise exposure by business boundary and entitlement type. It helps identify whether a variance is caused by population, rate, consumption, or allocation changes and gives organization owners a meaningful planning target.

Use a provider-enforced organization budget only where that capability and scope are confirmed. Otherwise, label the value as an internal allocation or monitoring threshold. Do not confuse it with an organization policy profile, and do not assume organization totals sum cleanly to enterprise totals when shared users, direct enterprise assignments, delayed records, or exclusions exist.

13.3.1 Copilot Business

Configure the Copilot Business budget for AI credits.

Define the Copilot Business allocation from the current number of eligible Business seats, the provider's applicable included-credit rate, observed usage, planned population changes, and approved overage posture. Link the allocation to the organizations, cost centers, and owners responsible for those seats.

Separating this SKU supports accurate entitlement and consumption forecasting and prevents a blended enterprise average from hiding a Business-specific variance or adoption pattern.

Do not reuse rates, features, or assumptions from another Copilot SKU. Seat counts and assignments can change during the period, and a billed seat record may not match current membership at every point in time. Preserve the rate basis and effective dates, and reconcile provider totals before changing the budget.

Budgets and alerts / Edit Copilot Business

Edit Copilot Business

SKU

Copilot Business

Budget scope

Organization

Scoped budgets track spending for all repositories in the organization

 autocloudarc-digital-services

Budget

Set a budget amount to track your spending on a monthly basis

License count

10 ~ \$190 budget amount

\$19/license per month for Copilot Business

☐ Stop usage when budget limit is reached
Not available for license-based products

Alerts

Receive emails when usage reaches 75%, 90%, and 100% of the budget amount

☒ Receive budget threshold alerts

Alert recipients

Account owners and billing managers receive budget alert emails by default.
You can add other recipients if needed.

Select recipients ▾

1 selected

 autocloudarc

×

Update budget

Cancel

Figure: 8 - Copilot Business organization budget configuration for autocloudarc-digital-services.

13.3.2 Copilot Enterprise

Configure the Copilot Enterprise budget for AI credits.

Define the Copilot Enterprise allocation from the current eligible Enterprise seat inventory, applicable included-credit rate, supported capabilities, observed consumption, and forecast demand. Record the responsible organizations, cost centers, owners, and any material workload expected to affect usage.

This separate allocation makes the financial effect of the Enterprise SKU and its usage patterns visible. It also supports informed decisions about license mix, feature posture, training, and exceptions without distorting the Business baseline.

A higher entitlement or broader feature set does not prove that every user needs a higher ULB or that all consumption creates equal value. Rates and eligible features can change, and cross-organization or enterprise-direct assignments may complicate attribution. Validate the current basis and evaluate outcomes as well as volume.

14. Monitor Usage

Examine enterprise GitHub Copilot usage insights for trends, adoption, and potential areas for optimization.

Usage monitoring is the repeated observation of AI Credit consumption, burn rate, remaining entitlement, forecast exposure, alerts, data freshness, license inventory, and profile assignments. Define the source, owner, frequency, thresholds, expected response, and escalation path for every monitored signal.

Monitoring provides early warning before a financial or user-level boundary is reached. It also reveals failed data collection, unexpected concentration, assignment drift, and changes in adoption that require investigation.

Usage volume alone does not establish business value, misuse, or waste. Provider data can be delayed, revised, aggregated, or incomplete, and a dashboard trend may reflect a population or entitlement change rather than behavior. Display freshness and coverage, preserve unknown values, and investigate causes before changing a profile.

14.1 Weekly

Review weekly GitHub Copilot usage insights for trends, adoption, and potential areas for optimization.

The weekly review is an operational check of active alerts, rapidly changing burn rates, high-exposure cohorts, pilots, exceptions, incidents, failed data sources, and assignments due to expire. Record actions, owners, acknowledgement status, and decisions that cannot wait for period close.

This cadence leaves time to validate an emerging issue, communicate with affected users, and choose an approved response before a monthly boundary is reached. It is especially useful during rollout or when forecast uncertainty is high.

Weekly data is a partial-period signal and can be distorted by working-day patterns, releases, leave, or one-time workloads. Do not annualize a short spike or penalize a cohort from an isolated week. Compare like periods and use the review to trigger investigation rather than to make automatic long-term changes.

14.2 Monthly

Review monthly GitHub Copilot usage insights for trends, adoption, and potential areas for optimization.

The monthly review reconciles the completed or substantially complete billing period across provider usage, included credits, license inventory, metered charges, budgets, cost centers, organizations, and profile assignments. Compare actuals with forecast and explain material variances.

Monthly evidence supports budget decisions, chargeback or showback, profile reassessment, trend analysis, and executive reporting. A consistent close process creates comparable baselines across periods.

Billing, finance, and calendar months may close at different times, and late events or provider adjustments can revise an apparent final value. Define a data cutoff and restatement process, retain the original snapshot, and avoid declaring the period reconciled while material sources remain incomplete.

15. Analyze Monthly AI Credits Usage

Review monthly AI credits usage for trends, adoption, and potential areas for optimization.

Monthly analysis converts reconciled usage into explanations and decisions. Evaluate total and per-user consumption, concentration, SKU and organization mix, forecast variance, overage, alert history, profile distribution, exceptions, and changes from comparable prior periods.

The purpose is to identify which factors drive exposure and whether controls are supporting approved outcomes.

Segmenting by population and control profile helps separate broad adoption from a small number of specialized or anomalous workloads.

Averages can hide concentration, and correlation does not prove that a policy, training event, or feature caused a change.

Distinguish measured, calculated, and forecast values; account for population and rate changes; and pair financial analysis with operational and qualitative evidence before recommending action.

16. Obtain Qualitative User Feedback

Collect qualitative feedback from users regarding their experience with AI credits and GitHub Copilot.

Qualitative feedback captures user outcomes and constraints that billing data cannot show. Use interviews, surveys, support themes, pilot reviews, or structured retrospectives to ask about task value, interruption, model and feature fit, workarounds, latency, training needs, and the effect of current limits.

This evidence helps distinguish valuable high consumption from avoidable use and reveals controls that look effective financially but block approved work. It also provides context for policy, budget, support, and training decisions.

Feedback is subject to selection, recall, and response bias and may contain sensitive code, business, or personnel information. Use approved collection and retention methods, avoid requesting unnecessary content, include affected low-usage and interrupted users, and do not treat a small set of opinions as a replacement for reconciled quantitative evidence.

17. Review and Adjust Policy Sets

Review and adjust GitHub Copilot policy sets based on usage insights and qualitative feedback.

Review policy profiles by comparing their approved objectives with current provider capability, effective settings, usage evidence, incidents, exceptions, support findings, and user outcomes. Retain, revise, split, consolidate, or retire a profile only through the documented approval and change process.

Periodic adjustment keeps access posture aligned with business need and changing risk. It also removes stale exceptions and settings that no longer produce the intended outcome.

Do not change policy settings merely to reduce a financial variance when a budget control is the correct mechanism. Avoid changing several settings and budgets at once because the effects become difficult to attribute. Version the proposal, test representative use cases, stage rollout, and preserve rollback evidence.

18. Review and Adjust Budgets

Review and adjust enterprise, user level, and organizational budgets based on usage insights and qualitative feedback.

Review each budget layer against reconciled actual, forecast accuracy, business outcomes, interruption evidence, exception use, population changes, rates, and risk tolerance. Adjust the enterprise overage boundary, universal ULB, cost-center overrides, or organization allocations independently according to the decision each control serves.

This review keeps financial exposure intentional while allowing approved work to continue. It can correct a limit that is routinely ineffective, unnecessarily restrictive, or based on an obsolete entitlement or population assumption.

Increasing a budget is not proof of value, and lowering one is not optimization. Do not chase a single anomalous month, normalize recurring overages without approval, or use a larger ULB as a substitute for resolving inefficient behavior. Model the change, assess aggregate exposure and interruption, approve exceptions, and monitor the next period against explicit success criteria.

19. Token Optimization Training

Provide training on token optimization strategies to maximize the efficiency of AI credits usage.

Token optimization training teaches users to obtain the required outcome with appropriate models, features, context, prompts, iteration patterns, and stopping criteria. Tailor examples to approved workflows and show when a larger context or more capable model is justified rather than promoting minimum consumption in all cases.

The goal is effective use: reduce avoidable repetition and unnecessary context while preserving quality, security, accessibility, and developer productivity. Training can also reduce accidental high usage caused by unclear requests, unsupported workflows, or misunderstanding of model and agent behavior.

Tokens and AI Credits are not necessarily equivalent units, and conversion can vary by provider, model, feature, or time. Do not promise a fixed saving without current evidence. Never encourage users to omit necessary context, bypass policy, expose sensitive data, or accept lower-quality results solely to reduce measured consumption; evaluate outcomes as well as credit volume.

20. Continuous Improvement and Iteration (Sections 14-19)

Loop continuously to review and iterate on the processes, policies, and budgets related to AI credits and GitHub Copilot usage to drive ongoing improvement and optimization.

Continuous improvement is the governed cycle that connects monitoring, monthly analysis, user feedback, policy review, budget review, and training. At each cycle, document the evidence baseline, hypothesis, approved action, owner, success measure, observation period, result, and decision to retain, revise, or roll back.

The cycle keeps controls current as provider capabilities, rates, populations, business needs, and risks change. It also turns isolated adjustments into an auditable learning process and carries unresolved findings into the next review with named owners.

Continuous does not mean constant configuration change. Frequent overlapping changes create fatigue, obscure causality, and weaken baselines. Prioritize material issues, allow enough observation time, preserve stable control periods, and balance cost with value, access, security, reliability, and user impact.

Appendix A: Design and assignment criteria

Review usage evidence for explicit business decisions. Do not assign a profile solely because a user or group consumed more than an average.

Criterion	Decision question
Usage evidence	Are sources current, complete enough, owned, and reconciled?
Business need	What approved outcome requires this control or access posture?
Forecast exposure	What consumption and financial range is plausible during the billing period?
Interruption tolerance	What happens if alerts fire or usage is limited?
Security	Are data, feature, service, and integration risks accepted?
License eligibility	Is each person eligible and correctly entitled for the proposed profile?
Ownership	Who funds, approves, operates, supports, and reviews the assignments?
Duration	Are the assignments standing, temporary, pilot-only, or event-bound?
Review date	When must evidence, need and value be reassessed?

Reject or defer an assignment when material evidence is unknown, ownership is missing, license eligibility is unresolved, or interruption impact has not been accepted.

Appendix B: Governance

architecture and workflow

AI CREDITS GOVERNANCE

Owned evidence → controlled change → verified operation

APPENDIX B
WORKFLOW

OWNED INPUTS Confirm authority, freshness, scope, and coverage before use.

PROVIDER

Usage · billing · settings
Alerts · supported controls

FINANCE

Rates · charges · funding
Forecast · risk tolerance

IDENTITY & LICENSE

Users · teams · organizations
Eligibility · assignments · SKUs

BUSINESS & RISK

Need · value · security
Impact · interruption tolerance

EVIDENCE

RULE

Preserve source + timestamp + limitations. Mark unavailable or partial values **UNKNOWN**.
Reconcile provider, finance, license, organization, and cohort records before a material decision.

SEVEN GOVERNANCE GATES Each gate produces evidence and requires a named owner for unresolved action.

01 PREPARE

SERVICE · BUDGET
SECURITY · LICENSE

ACT

Set scope, roles, sources, handling, tests, success measures, and rollback.

RECORD Scope · role matrix · source plan · test plan

EXIT

Owners accept authority, evidence handling, and decision rights.

02 BASELINE

SERVICE · BUDGET
LICENSE · ADMIN

ACT

Collect records, preserve gaps, reconcile totals, and explain variance.

RECORD Source register · worksheet · reconciliation

EXIT

Values reconcile, or limitations are formally accepted.

03 DESIGN

SERVICE · BUSINESS
BUDGET · SECURITY
LICENSE

FOUR DISTINCT FINANCIAL CONTROL LAYERS

ENTERPRISE SPEND UNIVERSAL ULB COST-CENTER ULB ORG / SKU ALLOCATION

KEEP SEPARATE Access-policy profiles · assignments · alerts · exceptions

RECORD Draft profile register · impact assessment

EXIT

Every proposal has evidence, rationale, owner, duration, and review date.

04 APPROVE

BUDGET · SECURITY
LICENSE · BUSINESS
SPONSOR

DECIDE

Profile · assignment · exception · pilot · support path · change record
RECORD Decision · approvers · impact · authorized change · rollback trigger

APPROVED FOR PILOT?

YES → limited change

NO → return / defer with owner

05 PILOT

PILOT LEAD · ADMIN
SERVICE · SUPPORT

ACT

Capture prior state, apply limited change, test actual behavior, monitor, reconcile.
RECORD Before / after state · tests · measures · feedback

SUCCESS + TOLERANCE?

YES → staged rollout

NO → restore prior state

06 ROLLOUT

SERVICE · ADMIN
CHANGE · BUSINESS

ACT

Sequence waves, reconfirm evidence, communicate, apply, verify, and pause.
RECORD Wave log · verification · communications · active register

EXIT

Assignments match the register; ownership is clear.

07 OPERATE

SERVICE · BUDGET
SECURITY · LICENSE
BUSINESS

ACT

Review, reconcile, resolve alerts, control changes, and retire stale profiles.
RECORD Reviews · incidents · changes · exceptions · trends

EXIT

Verified provider state matches the active register.

MONITOR & RESPOND

Display freshness and coverage. An alert is not proof that usage stopped.

75% · EARLY WARNING

Validate data → refresh forecast
Notify owners → confirm need

90% · CRITICAL

Escalate → assess continuity
Approve action → communicate

100% · LIMIT REACHED

Use tested enforcement, incident, exception, or rollback path

CADENCE Weekly · monthly · quarterly · event-driven

REVIEW DECISION

Use reconciled quantitative evidence + user outcomes + current provider behavior.

RETAIN → OPERATE

REVISE → DESIGN

ROLL BACK → BASELINE

Rollback restores captured provider state + prior active register and preserves history.

CLOSE ONLY WHEN VERIFIED PROVIDER STATE AND THE ACTIVE REGISTER AGREE.

ROLLBACK / RE-BASELINE BEFORE REPLY

Figure: 9 - AI Credits governance architecture and workflow.

Appendix C: Lifecycle phases

Each phase is a gate. Record a named owner for every unresolved action before moving forward.

Phase	Owners	Actions	Evidence	Exit criteria
Prepare	Service, budget, security, and license	Set scope, roles, sources, handling, tests, success measures, and rollback	Scope, role matrix, source plan, and test plan	Owners accept scope, authority, evidence handling, and decision rights
Baseline	Service, budget, license, and administrator	Collect sources, mark gaps, reconcile totals, and document variances	Source register, worksheet, and reconciliation record	Required values are reconciled or limitations are formally accepted
Design	Service, business, budget, security, license	Propose four-layer profiles, assignments, responses, names, and reviews	Draft profile register and impact assessment	Every proposal has evidence, rationale, owner, duration, and review
Approve	Budget, security, license, business, sponsor	Decide each profile, exception, pilot, support path, and change record	Decisions, approvals, exceptions, and authorized change	Each proposal is approved, rejected, or returned with a named owner
Pilot	Pilot lead, administrator, service, support	Capture prior state, apply a limited change, test, monitor, and reconcile	Before and after state, tests, measures, and feedback	Success and tolerance criteria are met, or rollback is completed
Roll out	Service, administrator, change, and business	Sequence waves, confirm evidence, communicate, apply, verify, and pause	Wave records, verification, communications, and register	In-scope assignments match the active register and ownership is clear
Operate	Service, budget, security, license, business	Review, reconcile, resolve alerts, control changes, and retire profiles	Reviews, incidents, changes, exceptions, and trends	Controls operate on schedule and verified state matches the register

Appendix D: Monitoring and response

Set triggers from risk tolerance, forecast uncertainty, lead time, and validated model or harness provider behavior. Do not assume that an alert blocks usage or that a displayed limit is enforced.

Stage	Configurable trigger	Response	Evidence
Early warning (75%)	Customer-selected signal that leaves response lead time	Validate data, refresh forecast, notify owners, and confirm need	Alert, owner acknowledgement, data check, and new forecast
Critical (90%)	Customer-selected signal requiring an owner decision	Escalate, assess continuity, choose action, and communicate	Decision, impact assessment, approval, and communication
Limit reached (100%)	Verified boundary event or equivalent provider state	Follow tested enforcement, incident, exception, or rollback path	Provider event, action record, approval, and reconciliation

For each stage, record trigger logic, source, evaluation frequency, recipients, acknowledgement target, decision authority, supported actions, and fallback communication. Validate alert timing and delivery with controlled tests.

Appendix E: Reconciliation and data quality

Use this procedure whenever sources disagree or coverage changes:

1. Preserve the original records, timestamps, scope, and warnings.
2. Classify each value by source and distinguish measured, calculated, forecast, and manually confirmed values.
3. Confirm period, account scope, population, service, entitlement, and rate basis from authoritative records.
4. Compare provider totals with finance, license, organization, and cohort records without treating unlike measures as interchangeable.
5. Separate expected coverage differences from unexplained variance.
6. Record the variance, cause, owner, tolerance decision, and due date.
7. Recalculate forecasts and reassess assignments with reconciled inputs.
8. Block approval when unknown data or information gaps could change eligibility, exposure, precedence, interruption impact, or enforcement decisions.

Treat failed, unauthorized, unsupported, delayed, hidden, or partial results as unknown. Never convert them to zero, an empty population, or evidence of no consumption.

Appendix F: Exceptions, change control, and rollback

Every exception must record the affected profile and population, business reason, risk, compensating controls, forecast impact, approvers, start, expiry, review date, and rollback trigger.

Use this change procedure:

1. Link the request to the active profile version and evidence baseline.
2. Capture current provider state and a tested restoration method.
3. Assess financial, security, entitlement, business, support, and interruption effects.
4. Obtain required approvals and create the proposed register version.
5. Validate through a controlled test or approved pilot.
6. Have an authorized administrator apply the approved change.
7. Independently verify scope, precedence, alerts, and enforcement behavior.
8. Observe for the approved period and reconcile results.
9. Close the change only after evidence and the active register agree.

Roll back when assignment is wrong, access is unauthorized, interruption is unexpected, variance exceeds tolerance, evidence is incomplete, or approved success criteria fail. Restore the

captured prior provider state and prior active register version.
Preserve decision, incident, test, and reconciliation history.

Appendix G: Review cadence

Choose the cadence that matches exposure and operating maturity. More than one cadence can apply.

Option	Review focus
Weekly	Active alerts, pilots, high-exposure cohorts, incidents, and exceptions
Monthly	Entitlements, consumption, cost, forecast, assignments, and reconciliation
Quarterly	Profile design, access posture, ownership, evidence, training, and drift
Event-driven	Rate, entitlement, feature, organization, risk, incident, or control change

At each review, confirm that source availability, rates, supported controls, scope, precedence, alert behavior, and enforcement assumptions remain current.

Appendix H: Getting-started checklist

1. Name the service, budget, security, license, business, and operating owners.
2. Define the first account, population, service, and reporting period.
3. Create the evidence source register and baseline worksheet.
4. Reconcile entitlement, consumption, cost, and existing control coverage.
5. Document provider-validation questions and controlled tests.
6. Draft one profile for each needed governance layer.
7. Approve a limited pilot, success measures, communications, and rollback.
8. Run the pilot, reconcile results, and decide whether to revise or roll out.
9. Start the selected review cadence and keep the active register current.

Appendix I: Evidence checklist

Before approving a profile or closing a change, confirm the evidence set has:

- Defined scope, owners, period, exclusions, and authorization context
- Current provider documentation and authoritative records used for decisions
- Completed source register with limitations and retrieval timestamps
- Reconciled baseline and approved variance treatment
- Profile version, assignment rationale, values, owners, and review dates
- Business, budget, security, license, and operating approvals as applicable
- Controlled test results for scope, precedence, alerts, and enforcement
- Exception records with compensating controls and expiry
- Prior state, approved change, independent verification, and observation record
- Communications, incidents, rollback actions, and post-change reconciliation

Appendix J: Glossary

Term	Definition
AI Credit (AIC)	Provider-defined unit for eligible AI consumption
User-Level Budget (ULB)	Customer-defined per-user boundary where supported
Financial guardrail	Approved control intended to bound aggregate financial exposure
Cohort override	Time-bound or standing profile for an approved higher-usage population
Access-policy archetype	Reusable access posture managed separately from financial controls
Profile register	Controlled record of profile versions, assignments, owners, and evidence
Reconciliation	Comparison that explains differences among authoritative and derived records
Unknown	Value that cannot be established from a complete, authorized source
Controlled test	Limited validation of actual provider behavior under approved conditions
Rollback	Restoration of the captured prior provider state and active profile version