

Enterprise AI Credits Governance

A phase-led operating model for budget controls, governance, and evidence

A practical guide to included AI-credit pools, per-user limits, metered budgets, controlled rollout, monitoring, and accountable review decisions.

OPERATING PRINCIPLE

Use the seven-phase governance model to move from evidence to design, approval, pilot, rollout, operation, monitoring, and review. Treat every production change as a typed Active Register record with a named owner, evidence, an effective period, and a tested rollback path.

PUBLICATION DISCLAIMER

Independent analysis and opinions only; not necessarily the views, positions, policies, or endorsements of the author's employer, affiliated organizations, GitHub, Microsoft, or providers. Based on official cloud-service and developer-platform documentation available at publication, and empirical examples from the author's GitHub Enterprise Cloud development account. Capabilities, pricing, policies, and availability may change; verify current provider information before technical, operational, financial, legal, or procurement decisions.

Provided "as is," without warranties (accuracy, merchantability, fitness, non-infringement, or availability); no liability, professional-services relationship, commitment, endorsement, or guarantee is created, to the maximum extent permitted by law.

Preston K. Parsard

Contributor: AI

Last updated: September 15, 2026

Confirm product capabilities, billing behaviour, and administrative roles in current GitHub documentation before production use.

Contents

Section	Purpose
Read first	Audience, roles, architecture, financial-control model, and the governance lifecycle.
01 Prepare	Set scope, owners, evidence sources, and the included-credit baseline.
02 Baseline	Assess current provider state and establish the typed Active Register.
03 Design	Define ULBs, budgets, naming, profiles, units, ownership, and tests.
04 Approve	Authorize a bounded change with evidence and a rollback decision owner.
05 Pilot	Verify scope, precedence, enforcement, alerting, and rollback.
06 Rollout	Apply profiles, ULBs, scoped metered budgets, and the enterprise boundary.
07 Operate	Run triage, reconciliation, analysis, learning, and the continuous Monitor & Respond / Review Decision overlay.
Appendix A	Design and assignment criteria.
Appendix B	Control decision sequence.
Appendix C	Lifecycle and evidence checklist.
Appendix D	Monitoring and response.
Appendix E	Reconciliation and data quality.
Appendix F	Exceptions, change control, and rollback.
Appendix G	Review cadence.
Appendix H	Getting-started checklist.
Appendix I	Governance glossary.
Appendix J	Active Register schema and operating guide.
References	Official GitHub documentation and FinOps source basis.

Read first

This guide is written for enterprise administrators, budget owners, service owners, security owners, finance partners, product and procurement partners, FinOps practitioners, and assurance reviewers who need one readable method for governing GitHub Copilot AI Credits. It separates **per-user AI-credit limits** from **aggregate metered-spend budgets**, and it keeps policy-profile configuration distinct from financial control. [1] [2]

AI FinOps scope

Use this guide as a GitHub Enterprise Cloud control implementation within a broader enterprise AI FinOps practice. The FinOps Foundation frames AI spending as granular, fast-moving, and often cross-category; this guide does not attempt to govern every AI vendor or cost type. [7]

How to use the guide

Follow the phases in order. Use the numbered procedures to configure or validate a control. Use the appendices as supporting reference, not as a substitute for phase exit gates. The Active Register is the single authoritative record for profiles, User-Level Budgets (ULBs), metered budgets, rollout waves, tests, and exceptions.

Acceptance-test quick reference

P-01 verifies User-Level Budget (ULB) precedence; P-01a verifies a named policy's effective behavior; P-02 verifies shared-pool and included-usage behavior; P-03 verifies ULB hard-stop behavior; P-04 verifies metered-budget alerting and Stop usage behavior; P-05 verifies rollback. Phase 05 defines the full test methods.

Where to host the Active Register

Use a controlled system with role-based access, change history, immutable evidence links, and export capability. An ITSM/GRC record or repository-backed YAML/JSON is preferred. A governed spreadsheet is acceptable only when it has protected key columns, controlled picklists, mandatory-field validation, restricted edit rights, version history, link validation, and an exportable phase-exit snapshot. Record the system of record, record owner, and workspace location in the Prepare evidence plan.

Audience and outcomes

By the end of the guide, the organization will have a reconciled entitlement baseline, documented control model, approved target state, pilot evidence, controlled production rollout, and a recurring monitoring and review cadence.

The guide uses a running principle: first identify the applicable per-user ULB, then the shared included-credit pool, then the relevant aggregate metered-spend guardrail. A configuration must never be assumed to enforce because it appears in a dashboard or screenshot. [1] [2]

Roles and prerequisites

Role	Accountability
Executive sponsor	Accepts risk posture and material unresolved exceptions.
Enterprise budget owner	Approves enterprise financial exposure, threshold posture, and escalations; typically an authorized enterprise owner or billing manager.
Service owner	Owns readiness, evidence, communications, and operational health.
Security owner	Approves security posture, policy-profile constraints, and compensating controls.
License owner	Validates entitlement inventory, eligible population, and license assignment; maps to the GitHub licensing or billing access needed for the selected scope.
Cost center owner	Accepts cost-center membership, funding, and review obligations.
Authorized administrator	An enterprise owner, billing manager, or appropriately authorized organization owner who applies approved provider changes and verifies configured state.
Finance / FinOps partner	Supports attributable reporting, forecast review, funding decisions, and value-aware governance; may coordinate an AI Investment Council or equivalent review forum. [7]
Product / portfolio owner	States the investment stage, intended outcome, and value hypothesis for material AI-enabled work. [7]
Procurement partner	Assesses acquisition channels, vendor commitments, and rate options when those decisions affect the approved scope. [7]
Assurance reviewer	Independently reviews evidence, reconciliation, and active-register integrity.

Prerequisite

Name the people—not just role titles—who approve, operate, receive alerts, support users, and decide on rollback. Capture their delegates and escalation route in the Active Register.

Optional investment forum

For material or rapidly changing AI initiatives, use an AI Investment Council or equivalent review forum to connect engineering, finance, product, procurement, and leadership decisions. This is an operating recommendation, not a GitHub product requirement. [7]

Architecture and workflow

Governance operating model

Scope of this section

This is a governance workflow and decision-path architecture, not a platform-topology or systems-integration diagram. It shows how administrators sequence controls, evidence, and decisions.

The first view explains the financial decision path. The second view turns that path into a governed lifecycle with entry gates, evidence, and review decisions. Use both views together: a budget setting is not complete until its scope, unit, owner, evidence, enforcement behaviour, and rollback path are documented.


Decision-path walkthrough

Step	Reader instruction
1. User boundary	Evaluate individual, then cost-center, then universal ULB. If the applicable ULB is exhausted, the request is blocked. [2]
2. Shared entitlement	If the ULB has headroom, usage draws from the shared included-credit pool while capacity remains; capture provider evidence rather than manually tallying usage. [1]
3. Cost-center / organization route	If a cost-center included-usage cap is reached, record the configured behavior: block members at the cap or continue their additional usage as paid overage. This control has no aggregate-budget Stop usage toggle. When paid overage is permitted, determine the applicable cost-center or organization budget and its allocation evidence; see Section 3.4 and Appendix J.4. [2] [5]
4. Enterprise and decision loop	If no narrower budget applies, evaluate the Enterprise spending limit. Record the signal and make a retain, revise, or rollback decision through the governance loop. [2]

How to use the visual

The detailed reference visual on the following page is a map. Read the four steps above first, then use the visual to locate the same decision points and evidence nodes. The written walkthrough is the accessible operational instruction.

AI-credit budget and consumption decision path



Use the written walkthrough for detailed scope, evidence, alerting, and rollback instructions.

Figure 1. Simplified AI-credit budget and consumption decision path. Use the written decision-path walkthrough on the preceding page as the accessible operational instruction.

Governance operating model

Recommended control-lifecycle gates




Figure 2. AI Credits Governance operating model: seven lifecycle phases supported by a continuous monitoring and decision loop.

Sequence rule

A team moves through the seven lifecycle phases—Prepare, Baseline, Design, Approve, Pilot, Rollout, and Operate. Monitor & Respond and Review Decision are continuous governance-loop nodes: they may be triggered during any phase and are sustained through Operate. They are not additional linear phases.

Financial control and ownership model

Five control layers

Layer	Scope / unit	Enforcement and owner
1. User-Level Budget (ULB)	Per user; AI Credits; individual > cost center > universal.	Hard stop for supported AI-credit-consuming surfaces. Universal: enterprise budget owner approves. Cost center: cost center owner approves. Individual: enterprise budget owner approves. Named service owner supports. [2]
2. Shared included-credit pool	Billing entity; AI Credits; monthly pooled entitlement.	Not a per-user assignment budget. License owner reconciles. [1]
3. Cost-center included-usage control	Cost center; AI Credits; before the metered phase.	Provider-calculated cap based on assigned licenses. At the cap, record whether members are blocked or continue as paid overage; only the paid-overage route can use a downstream metered-budget record. It has no aggregate-budget Stop usage state. [2] [5]
4. Scoped metered budget	Cost center or organization; US dollars (USD); metered phase only.	Alerts by default; a hard stop only if Stop usage is enabled. Cost center owner or organization owner operates. [2]
5. Enterprise spending limit	Enterprise; US dollars (USD); metered phase only.	Final enterprise-wide paid-usage boundary; not total bill. Enterprise budget owner operates. [2]

Layered-guardrail rule

Treat user, cost-center, organization, and enterprise controls as concurrent guardrails rather than a top-down allocation hierarchy. Use the included-usage phase to protect fair access to the shared pool, then use the metered phase to govern approved additional spend. A ULB spans both phases; aggregate metered budgets do not. [6]

Control precedence and units

Use **AI Credits** for ULBs and shared entitlement. GitHub documents aggregate additional-usage budgets in **US dollars (USD)**; one AI Credit equals \$0.01 USD. Treat the dollar value as the provider budget unit and keep the AI-credit value as consumption evidence. Do not create an internal alternate-currency enforcement rule. [1] [2]

Scenario	Controlling action
User reaches ULB	Block the request; verify individual > cost-center > universal assignment, remaining amount, and reset boundary.
ULB has headroom; pool remains	GitHub provider metering serves the request from the shared pool. Capture a provider evidence snapshot; do not manually tally usage.
Pool exhausted; cost center applies	Evaluate cost-center metered budget; block only if Stop usage is enabled and the budget is reached.
Pool exhausted; no cost center applies	Evaluate the relevant organization budget, then enterprise spending limit as applicable.
Several controls apply	Follow GitHub's documented flow: ULB, shared pool, then applicable cost-center, organization, or enterprise budget. The lowest remaining headroom can block first; preserve allocation/exclusion evidence and use P-04 to test the approved tenant scenario. [2]

Aggregate-budget routing after pool exhaustion

Observed allocation path	Admin design and test rule
User directly assigned to a cost center	Plan and test the named cost-center included-usage control before the metered phase and the cost-center metered budget after pool exhaustion; Stop usage applies to the latter. [2] [5]
No direct cost center; license billed to an organization	Plan and test the organization metered budget for the billing organization. For multi-organization licensing, document the provider allocation evidence before relying on the organization budget. [2]
No applicable cost-center or organization budget	Plan and test the Enterprise spending limit as the enterprise-wide post-pool USD boundary. [2]
Several controls appear relevant at once	Use the documented flow: ULB, shared pool, then cost center, organization, or enterprise metered budget. Capture cost-center exclusion and allocation evidence, then use P-04 to validate the approved tenant scenario and any lower-headroom interaction. [2]

Concurrent-control rule

GitHub documents the request flow and a lowest-remaining-headroom effect across applicable controls. Use P-04 to preserve tenant-specific allocation, cost-center-exclusion, alert, and interruption evidence; do not use a pilot to override published control semantics. The enterprise budget owner resolves any material evidence conflict before production change. [2]

Ownership and control matrix

Control record	Approving role	Operating owner	Rollback decision owner
Enterprise spending limit	Enterprise budget owner	Enterprise budget owner	Executive sponsor
Organization metered budget	Organization owner	Organization owner	Enterprise budget owner
Cost-center included-usage control	Cost center owner	Cost center owner	Enterprise budget owner
Cost-center metered budget	Cost center owner	Cost center owner	Enterprise budget owner
Universal ULB	Enterprise budget owner	Named service owner	Executive sponsor

Control record	Approving role	Operating owner	Rollback decision owner
Cost-center ULB	Cost center owner	Cost center owner	Cost center owner
Individual ULB	Enterprise budget owner	Named assigned owner	Service owner
Policy profile / rollout wave	Security + service owner	Authorized administrator	Executive sponsor

Role-to-control and configuration-surface map

Control	Apply surface and guide role	Verify and approve
Universal / individual ULB	Enterprise budget-control surface; enterprise budget owner directs an enterprise owner or billing manager.	Verify effective user record; complete ULB precedence and hard-stop tests; enterprise budget owner approves.
Cost-center ULB / included-usage / metered budget	Selected cost-center surface; cost center owner directs an enterprise owner or billing manager with confirmed cost-center privilege.	Verify direct-user allocation when relevant; complete ULB, P-02, or P-04 test; cost center owner approves.
Organization metered budget	Organization billing or budget surface; organization owner directs the provider role confirmed for that organization.	Verify billing scope and metered-budget test; organization owner approves.
Enterprise spending limit	Enterprise billing or budget surface; enterprise owner or billing manager applies it.	Verify US-dollar amount, scope, and metered-budget test; enterprise budget owner approves; executive sponsor accepts only exceptional exposure or rollback decisions.
Policy profile	Enterprise or organization Copilot policy settings surface; security owner directs an authorized administrator.	Verify configured and effective setting; security and service owners approve.

Provider access check

GitHub roles and UI labels vary by account and scope. Before production change, confirm the current provider path and privileges in GitHub documentation. This guide's roles define accountability; they do not replace provider authorization. [2] [3]

PHASE PURPOSE

Prepare establishes scope, authority, evidence sources, and the entitlement baseline. It does not apply production budgets, policy profiles, or ULB assignments.

WHAT THIS PHASE DELIVERS

- Named owners and delegates.
- Authoritative provider, license, finance, and evidence sources.
- Timestamped shared-pool entitlement baseline.
- Initial Active Register shell and evidence repository.

Exit gate

Scope, source responsibility, and evidence expectations are accepted; configuration remains in design.

1.1 Capture the included-credit baseline

Credits are pooled at the billing-entity level, unused credits do not carry over, and the shared included-credit pool resets at **00:00:00 UTC on the first day of each calendar month**. Capture actual provider records rather than estimates. [1]

For each billing entity, create an **entitlement-baseline** Active Register record. Record the immutable billing scope, provider source or evidence URL, retrieval timestamp, entitlement basis, included-credit amount when exposed, billing-cycle start/end, observed next-reset date, time-zone evidence, named owner, and **baseline_purpose** of forecast-only, reconciliation-only, or both. The documented pool-reset rule does not remove the need to preserve tenant evidence, especially when validating ULB reset behavior or reconciled reports.

- 1 Use the current enterprise Billing & licensing > AI usage or budget experience; labels can change, so confirm the current provider path and access role in References [1]–[3].
- 2 Record billing entity ID, organization or enterprise scope, SKU label or ID, eligible license count, entitlement basis, billing-cycle start/end, observed next reset, time zone if shown, source page URL, and retrieval timestamp.
- 3 Capture a provider evidence snapshot. Preserve unknown, delayed, excluded, or restated values; do not silently replace them with estimates.
- 4 Reconcile the reported pool to the licensed population. Record differences, source freshness, accountable owner, next review date, and Appendix E data-quality status.

PHASE PURPOSE

Baseline records what is configured and effective today before the target state is designed. It protects the organization from building controls on assumed provider behaviour.

WHAT THIS PHASE DELIVERS

- Current enterprise and organization policy inventory.
- Current budget, ULB, and assignment evidence.
- Attributable usage and cost evidence, or a documented source gap.
- Typed Active Register with source freshness and known gaps.
- Recorded exceptions and observed effective-state differences.

Exit gate

Current state, source freshness, attribution gaps, unknowns, and material variances are recorded and reviewed.

2.1 Assess current policy and budget state

Retrieve the current enterprise and organization configuration. For each policy, ULB, budget, and assignment, distinguish configured value from effective behaviour. Record inheritance, overrides, exceptions, provider limitations, and evidence source.

Do not assume equal configuration produces equal effective behaviour across scope. Validate a representative user, affected organization, and billing record during the controlled pilot.

Before optimizing budgets, model guidance, or policy profiles, create an attribution view that maps available AI usage and cost evidence to the applicable users, cost centers, organizations, and accountable owners. If a source cannot support attribution, record the gap, its decision impact, and the owner/due date rather than guessing. [6] [7]

Attribution before optimization

Use the available GitHub AI usage reporting and the organization's approved reporting pipeline to establish who consumed what and at what cost before right-sizing controls. Treat report freshness, allocation basis, and known omissions as evidence attributes. [6]

2.2 Establish the Active Register

The Active Register is a single typed record set. Record types are **policy-profile**, **ULB**, **entitlement-baseline**, **included-usage-control**, **metered-budget**, **license-baseline**, **rollout-wave**, **controlled-test**, and **exception**. The Profile Register is the filtered view of the Active Register's policy-profile records.

Field group	Required fields
Core identity	record_type; enterprise_control_id; scope_id; provider_control_id only when exposed (evidence/linkage, not key material). When no provider ID is exposed, capture source_url + retrieval_timestamp.
Ownership and support	owner_primary required from Prepare. owner_delegate and support_contact required before production approval. See Appendix J.2 for field definitions and lifecycle validation.
Conditional profile fields	policy_or_profile and profile_version only when a named control/profile pattern is used; otherwise use not-applicable / 0 in an export key suffix.
Assignment and period	assignment_target; effective_start/effective_end; approval_id. Timing fields (billing_cycle_start/billing_cycle_end; observed_next_reset; timezone_source) apply only where cycle/reset evidence is required: ULB, included-usage-control, and entitlement-baseline. See Appendix J.3/J.4.
Control values	• ULB: ulb_type; ulb_amount_ai_credits (whole AI Credits). • Entitlement baseline: entitlement_basis; included_credits_amount; baseline_purpose. • Metered budget: budget_amount_currency (numeric-only, no symbol/text); currency=USD; covered_ai_credit_sku (one or more labels/IDs); alert_thresholds (ordered provider-supported values); alert_recipients; response_sla; stop_usage_state. • License baseline: license_baseline_amount (numeric-only); license_baseline_currency (ISO code). • Included usage: downstream_metered_budget_reference only when applicable; never stop_usage_state. See Appendix J.3/J.4 for allowed values and record-type rules.
Evidence and exceptions	evidence_link; test_id/test_result; observed behavior at included-usage cap; exception_id/exception_expiry; support_contact; rollback_reference
Provenance	source_url; retrieval_timestamp; record_status; notes

Logical key

Use (**record_type**, **scope_id**, **enterprise_control_id**) as the base canonical logical key. Append (**policy_or_profile**, **profile_version**) only when the record uses a named profile or control pattern; otherwise represent the suffix as **not-applicable**, **0** in an export. **enterprise_control_id** is the stable internal identifier; **provider_control_id** is provider evidence, not key material. Owner name is a required attribute, not key material. Starting in Approve, require a rollback reference for every production-intended record; Prepare and Baseline records may use a planning status until a rollback path is designed.

Core identity and linkage example — metered-budget record

This is an identity/linkage excerpt, not a complete valid metered-budget record. Use the full worked example in Appendix J.5 and the mandatory-field matrix in Appendix J.4 before implementation.

Field	Example value
record_type	metered-budget

Field	Example value
enterprise_control_id	ENT-CTRL-0003
scope_id	ccr-0012-it
provider_control_id	budget-ccr-0012-it-ai-credit-v1 (when exposed)
policy_or_profile	ai-credit-guardrail
profile_version	v1
owner_primary	cost-center-owner
owner_delegate	finance-delegate
evidence_link	https://records.example/evidence/ccr-0012-it/2026-08-26
rollback_reference	CHG-1042; prior-state snapshot 2026-08-25

PHASE PURPOSE

Design defines the target state. It makes scope, units, ownership, enforcement intent, acceptance tests, and rollback requirements explicit before approval.

WHAT THIS PHASE DELIVERS

- ULB taxonomy, naming, and cohort mapping.
- Enterprise, cost-center, and organization metered-spend budget designs.
- Reusable enterprise and organization policy-profile designs.
- Typed Active Register records and a complete acceptance-test plan.

Exit gate

Every proposed control has scope, unit, owner, rationale, period, expected behaviour, evidence, test, and rollback path.

3.1 Design User-Level Budget profiles

Technical precedence. A User-Level Budget (ULB) is a per-user AI-credit boundary. Individual ULB overrides cost-center ULB; cost-center ULB overrides universal ULB. The applicable ULB is a hard stop. [2]

Operational design sequence. Draft the universal ULB first, then approved cost-center cohorts, then individual exceptions. Do not assign production users until taxonomy, membership validation, and pilot plan are approved.

Reset evidence. Do not infer a calendar-month or time-zone boundary for a ULB. Record tenant-observed next-reset evidence, the provider surface used, any exposed time-zone information, and the controlled test used to verify reset behaviour. The shared included-credit pool has a separately documented reset rule in Section 1.1; do not apply that rule to a ULB without tenant evidence.

Usage surface	Governing control and expected behaviour	Verification
AI-credit-consuming Copilot features	ULB; AI Credits; hard stop at the applicable ULB.	Confirm current product coverage and run P-03 on each intended surface. [2]
Code completion, next edit suggestions, or any named product surface	Coverage must be confirmed in current GitHub documentation; no ULB or metered-spend coverage is assumed without dated evidence.	Record the exact provider product label, coverage decision, source URL, and retrieval timestamp before rollout. [2]

Usage surface	Governing control and expected behaviour	Verification
Any unlisted AI product or surface	No coverage assumed.	Classify it in Design; record the coverage decision in the Active Register before rollout.

Coverage warning

ULBs and aggregate budget controls apply only to supported AI-credit-consuming usage. Do not assume they govern every Copilot feature, plan element, or other AI service. Before rollout, record the intended product surface, provider SKU label or ID, coverage decision, source URL, retrieval timestamp, and P-03 outcome in the Active Register.

Taxonomy legend

abd = abundant-usage cohort. The **k** suffix means whole thousands of AI Credits: 6k = 6,000. Add a cohort abbreviation only after it has an approved definition and a matching naming example; do not maintain unused abbreviations.

ULB type	Naming example	Design rule
Universal	user-universal-standard-6k	Default per-user amount for eligible users.
Cost-center cohort	user-cost-center-ccr-0012-abd-7k	abd = abundant-usage cohort; stable cost-center ID + whole-thousand AI-credit band.
Individual exception	user-individual-<user-id>-exception-8k	Documented exception, named owner, expiry, and review date.

3.2 Validate cost-center membership

Cost centers can contain users, repositories, organizations, and enterprise teams; for Copilot AI-credit allocation, membership can be direct, organization-derived, or enterprise-team-derived. Direct user assignment controls allocation first. If no direct user assignment applies, allocation can follow the organization that granted the user’s Copilot license. Record the supported membership type, allocation basis, current documentation or UI URL, and retrieval timestamp for the selected setup. [2] [5]

Blocking rule. Before pilot or rollout, generate a user-to-cost-center allocation report for the affected population. The report must show the provider allocation basis, zero unresolved stale direct assignments, and any organization-based fallback. An internal primary-cost-center field or HR/finance tie-breaker is governance metadata; it does not override provider allocation behavior. Record resolution and owner in the Active Register, and block the wave until it is complete.

3.3 Design policy profiles

A **policy profile** is an internal governance bundle, not necessarily a GitHub product object with that exact name. It maps to named GitHub enterprise or organization Copilot policy settings, feature surfaces, assignments, and documented effective behaviour. Record the actual provider settings and surfaces in the Active Register.

Use reusable, versioned posture names. Enterprise profile: **ent-ghcp-policy-set-<posture>-v<major>**. Organization profile: **org-ghcp-policy-set-<posture>-v<major>**. The organization assignment belongs in the Active Register, not in a reusable profile name.

Increment the major version whenever behaviour or scope changes materially enough to need new approval, pilot evidence, or rollback treatment. Retain the prior active version until rollout verification passes.

Record	Example	Rule
Enterprise profile	ent-ghcp-policy-set-standard-v1	Reusable enterprise posture; scope is ent.
Organization profile	org-ghcp-policy-set-standard-v1	Reusable organization posture; scope is org.
Actual assignment	standard-v1 → autocloudarc-digital-services	Assignment, owner, period, evidence, and rollback are register fields.

3.4 Design included-usage and metered-spend controls

A **cost-center included-usage control** is distinct from a budget: it caps a cost center's draw from included AI Credits before the metered phase and is calculated from assigned licenses. Aggregate cost-center, organization, and enterprise budgets activate only after the shared pool is exhausted, are denominated in **US dollars (USD)**, and do not extend a ULB. Stop usage is off by default for aggregate budgets; enable it for every approved hard-stop guardrail and verify effective behavior in P-04. [1] [2] [5]

Continuity-aware hard-stop decision

Use accountable owners and alert response as the primary guardrail. Approve Stop usage only when the accepted financial exposure of continuing usage is outweighed by the tested interruption risk to people, critical workflows, and business continuity. This is a governance decision; P-04 must still verify provider behavior for the named scope. [6]

Showback or chargeback decision

After attribution is established, choose whether to publish consumption to cost owners (**showback**) or apply an approved internal cost allocation (**chargeback**). Record the allocation basis, period, accountable owner, decision authority, and known limits; neither practice changes GitHub provider control behavior. [6] [7]

1	For a cost-center included-usage control, record the cost center, directly assigned license evidence, provider-calculated included-credit cap, configured/observed behavior at the cap (block or paid-overage), owner, evidence, and P-02 test. If paid overage is allowed, record the optional downstream_metered_budget_reference; do not assign the aggregate-budget Stop usage state to the included-usage control.
2	For the enterprise spending limit, record billing entity, US-dollar amount, period, covered_ai_credit_sku, Enterprise budget owner approval, alerts, Stop usage intent, emergency response, and P-04 test. Record executive sponsor acceptance only for exceptional exposure or a rollback decision.
3	For a cost-center metered budget, record cost-center scope, AI-credit SKU coverage, USD amount, recipients, Stop usage state, owner, evidence, and reconciliation plan.
4	For an organization metered budget, record billing organization, applicable AI-credit SKU, USD amount, thresholds, Stop usage state, provider support, owner, allocation caveat, and P-04 test.
5	For each control, distinguish an informational license-spend baseline from a pre-metered included-usage cap and an enforceable post-pool paid-usage guardrail.

Control	Scope / unit	When active and what to verify
Cost-center included-usage control	Cost center / AI Credits	Before the metered phase; provider-calculated from assigned licenses. At the cap, record block or paid-overage behavior. Only the paid-overage route may reference a separate downstream metered budget; that budget—not this control—holds Stop usage state. [2] [5]
Cost-center metered budget	Cost center / USD	After pool exhaustion; alerts by default; enable Stop usage for a hard stop; validate in P-04. [2]
Organization metered budget	Organization / USD	After pool exhaustion for licenses billed to that organization; enable Stop usage for a hard stop; validate allocation in P-04. [2]
Enterprise spending limit	Enterprise / USD	After pool exhaustion; alerts by default; enable Stop usage for a hard stop; validate in P-04. [2]
License-spend baseline	License population / forecast amount	Informational forecasting and allocation; not an AI-credit enforcement control.

3.4.1 Cost-center metered budget

A cost-center metered budget is an aggregate **US-dollar** guardrail after shared-pool exhaustion. It is not a per-user ULB and does not change a member's ULB. The configuration example below uses the `ccr-0012-it` cost center. [2]

SKU

All AI Credit SKUs

A budget for all SKUs that consume AI credits, where [overages spend is enabled](#).

Supported SKUs:

- Copilot AI credits
- Cloud Agent AI credits
- Spark AI credits

Budget scope

Cost center

Scoped budgets track spending for a single cost center

 ccr-0012-it

Budget

Set a budget amount to track your spending on a monthly basis

Budget amount

☒ Stop usage when budget limit is reached

Spending won't exceed your set budget

Alerts

Receive emails when usage reaches 75%, 90%, and 100% of the budget amount

☒ Receive budget threshold alerts

Figure 3. Cost-center metered budget configuration: cropped to the scope, budget, and enforcement panels for A5 readability.

Field	What to verify
Budget scope	One accountable cost center; confirm membership and owner.
SKU	Selected AI-credit SKUs; record the exact provider SKU label or ID, configuration-surface URL, retrieval timestamp, and covered_ai_credit_sku evidence in Appendix J.3.
Budget amount	Active Register value: budget_amount_currency=40 and currency=USD. A UI may display \$40, but the record stores the numeric amount and currency separately.
Stop usage	Enable only when an approved hard stop is required and supported.
Alerts	Record alert_recipients, response_sla, response owner, evidence source, and review date.

Contrast

A **cost-center ULB cohort** limits each member’s AI Credits. A **cost-center metered budget** limits the cost center’s aggregate US-dollar spend after pool exhaustion. Never use one control’s unit, owner, or enforcement assumption for the other.

3.4.2 Organization metered budget and license baseline

Forecasting and paid-usage control are deliberately separate records. Treat the license baseline (**record_type=license-baseline**) as an allocation and forecasting input, then design the organization metered budget (**record_type=metered-budget**) as the separate post-pool guardrail.
[1] [2]

3.4.2.a Forecasting: license-spend baseline — not an enforcement control

Edit Copilot Business

SKU

Copilot Business

Budget scope

Organization

Scoped budgets track spending for all repositories in the organization

 autcloudarc-digital-services

Budget

Set a budget amount to track your spending on a monthly basis

License count

10

~ \$190 budget amount

\$19/license per month for Copilot Business

☐ Stop usage when budget limit is reached

Not available for license-based products

Alerts

Receive emails when usage reaches 75%, 90%, and 100% of the budget amount

☒ Receive budget threshold alerts

Figure 4. Organization license-spend baseline configuration: cropped to the relevant forecast and scope panels; forecast only, not an enforcement control.

No Stop usage; informational only

A license-spend baseline forecasts eligible license charges. It neither limits AI Credits nor blocks paid usage. Record it separately from every ULB and metered-budget record.

3.4.2.b Enforcement: organization metered budget

Use case	Required record
License-spend baseline (record_type=license-baseline)	Organization, SKU, eligible licenses, monthly forecast amount, recipient, period, and variance owner.
Organization metered guardrail (record_type=metered-budget)	Billing organization, enumerated covered AI-credit SKU label(s) or ID(s) when available, USD budget amount, thresholds, alert_recipients, response_sla, Stop usage state, owner, allocation evidence, and P-04 test result. Retain any "All AI Credit SKUs" selector only as timestamped observed evidence.

Decision rule

Use a license baseline for forecasting and allocation. Use an organization metered budget for paid-usage control after pool exhaustion. Capture each as a separate typed Active Register record and follow Appendix B for routing.

3.4.2.1 Organization-scoped AI Credits budget example

The supplied configuration below fills the organization-scope gap between the forecast-only license baseline and the enterprise-wide spending limit. It is an **organization-scoped metered AI Credits budget for autocloudarc-digital-services**, not a per-user ULB and not a license-spend baseline. GitHub documents this control as capping metered charges for users whose Copilot licenses are billed to that organization after the shared pool is exhausted. The screen's repository wording is an observed configuration label, not a universal enforcement statement. [1] [2]




Figure 5. Organization-scoped AI Credits budget for autocloudarc-digital-services; supplied configuration screenshot.

Figure 5 configuration interpretation

Configuration item	What the screenshot shows and what to record
SKU coverage	"All AI Credit SKUs" is the literal observed UI selector. The canonical covered_ai_credit_sku field stores enumerated provider labels or IDs when available; store the selector text only as observed evidence or notes with source URL and retrieval timestamp.
Budget scope	Observed screen: Organization: autocloudarc-digital-services. Record the immutable organization or billing-scope ID. Documented behavior: metered charges for users whose Copilot licenses are billed to that organization; do not treat repository wording as universal scope.
Monthly amount	\$20 USD in the observed configuration. Treat this as an illustrative approved amount for the named organization and period, not a universal threshold.
Stop usage	Enabled in the supplied configuration. Before production use, record current support for the selected scope/SKUs, approval, evidence URL, and P-04 outcome; do not assume the posture applies beyond supported provider behavior.
Alerts and recipients	Observed in the supplied configuration: threshold alerts at 75%, 90%, and 100%; account owners and billing managers are shown as default recipients, with an additional named recipient selected. Verify actual defaults in current GitHub documentation and the configuration UI, then record primary owner, delegate, support contact, alert_recipients , and response_sla .

Active Register record — organization metered budget

Canonical field	Value / recording rule
record_type	metered-budget
scope_id	Immutable organization ID
policy_or_profile	organization-ai-credit-guardrail
profile_version	v1
budget_amount_currency	20
currency	USD
covered_ai_credit_sku	Enumerated provider label(s) or ID(s) when available; retain observed "All AI Credit SKUs" selector only as evidence or notes with source URL and retrieval timestamp.
alert_thresholds	75%;90%;100% — observed Figure 5 configuration; confirm current provider-supported values
stop_usage_state	enabled
owner_primary	Named accountable owner
owner_delegate	Named delegate
support_contact	Named support route
alert_recipients	Validated recipient set
response_sla	Approved response target
evidence_link	Current configuration and approval evidence
test_id	P-04
test_result	pass before production approval
rollback_reference	Approved prior-state restoration action
notes	Scope and allocation: this organization budget governs metered charges for licenses billed to this organization; cost-center and enterprise controls remain separate. Where a user is licensed by multiple organizations, GitHub may select one at random each billing cycle; record allocation evidence and prefer direct cost-center assignment when predictable enforcement is required. [2] [4]

3.4.3 Enterprise spending limit

The Enterprise spending limit is the enterprise-wide aggregate US-dollar control for paid AI-credit usage after shared-pool exhaustion. It does not cap license fees or act as a per-user ULB. [2]

Design field	Required decision
Scope and currency	Enterprise billing entity, provider billing cycle, US-dollar amount, covered AI-credit SKU, and effective period.
Financial posture	Approved budget amount, threshold bands, canonical alert_recipients, canonical response_sla, and exposure rationale.
Enforcement	Stop usage enabled for the approved hard-stop posture; capture current provider support and the P-04 evidence.
Response	Named primary owner, delegate, executive escalation, support route, and rollback or emergency-change authority.

Design exit check

The Enterprise spending limit has a separate typed record, not a substitute for ULBs, cost-center budgets, or organization budgets. Its pilot evidence covers configured state, effective state, alert behavior, Stop usage outcome, and provider allocation context.

PHASE PURPOSE

Approve turns a design into an accountable decision. No administrator applies a production control without a named authorizer, evidence, success criteria, and rollback decision owner.

WHAT THIS PHASE DELIVERS

- Signed or recorded approval decision.
- Complete Active Register record and evidence package.
- Pilot scope, test IDs, success criteria, and rollback decision owner.
- Communication and support plan for affected users.

Exit gate

The proposal is approved, rejected, or returned with a named owner and documented reason.

4.1 Create the approval record

- 1 Assemble scope, current baseline, proposal, owner, unit, expected behaviour, financial and security impact, support plan, and rollback method.
- 2 Link the decision to the typed Active Register record and capture approver, decision rationale, effective period, review date, and residual-risk acceptance.
- 3 Approve, reject, or return the proposal. A returned proposal must name the owner and missing evidence or redesign requirement.
- 4 Authorize a bounded pilot only after the test population, communication plan, success criteria, provider scope, and rollback path are explicit.

Approval exit check

The Active Register names the approver, operating owner, primary and delegate escalation contacts, support contact, rollback decision owner, approved scope, effective period, acceptance-test IDs, and required response SLA.

PHASE PURPOSE

Pilot validates the approved design in a limited scope. It records what the provider actually does, not what a configuration is expected to do.

WHAT THIS PHASE DELIVERS

- Representative user and organization or cost-center test scope.
- Configured-state and effective-state snapshots.
- Pass/fail evidence for precedence, pooling, alerting, enforcement, and rollback.
- Decision to roll out, revise, or roll back.

Exit gate

All planned tests have recorded results and any failure has an accountable decision.

5.1 Minimum acceptance-test suite

Test ID	Test	Pass condition
P-01	ULB precedence	Individual > cost center > universal assignment is effective for the approved test user.
P-01a	Policy precedence	For a named setting, configured enterprise and organization state is compared with effective behaviour; do not assume provider precedence without evidence.
P-02	Shared pool	Eligible request consumes shared entitlement while pool capacity remains.
P-03	ULB hard stop	Applicable supported ULB blocks the request on exhaustion at every intended AI-credit-consuming pilot surface; support impact is recorded.
P-04	Metered budget	During a non-critical window, a low approved US-dollar threshold produces timestamped alert evidence and expected Stop usage outcome for the named scope and period. Include concurrent-scope evidence or document why a combined-scope test is not possible.
P-05	Rollback smoke test	Prior provider state can be restored in the approved window and matches captured evidence.

Pilot test note. Use three to five approved users in a bounded window. Record the applicable ULB or low USD threshold, source timestamp, configured and effective state, alert or Stop usage outcome, support impact, and rollback evidence.

PHASE PURPOSE

Rollout applies only approved and piloted changes. Each wave is a typed Active Register record with scope, owner, dependencies, effective behaviour, evidence, and rollback criteria.

WHAT THIS PHASE DELIVERS

- Approved policy profiles applied and independently verified in controlled waves.
- User-Level Budgets assigned with verified precedence, reset evidence, and support readiness.
- Cost-center and organization metered budgets active with tested alert and Stop usage behaviour.
- Enterprise spending limit applied with reconciled scope, US-dollar amount, evidence, and rollback owner.

Exit gate

Configured and effective state agree with the Active Register; each completed wave has evidence and a named owner.

6.1 Apply policy profiles

Do not assume enterprise-versus-organization policy precedence. In the pilot, use P-01a to validate the effective behaviour of each named setting at the target scope. Then apply the approved enterprise or organization profile in controlled waves. Reconfirm profile version, population, owner, dependencies, communication plan, support readiness, and rollback condition. Pause the next wave if configured or effective state differs from the Active Register.

6.2 Apply User-Level Budgets

Where to configure and verify

Use the current GitHub Enterprise budget-control experience. An enterprise owner or billing manager applies universal and individual ULBs at enterprise scope; a cost-center owner directs an authorized administrator for cost-center ULBs. Confirm the exact available UI or API path, assigned scope, and effective user record against current GitHub budget documentation before production use. [2] [3]

- 1 Confirm the approved ULB record, population, AI-credit amount, named owner, effective period, reset verification, and rollback decision owner.
- 2 Apply universal ULB first; cost-center ULB only after membership validation; individual ULB only for an approved exception.
- 3 Verify configured and effective individual > cost-center > universal precedence, expected hard-stop behaviour, and the upcoming reset boundary for the test users.

-
- | | |
|---|---|
| 4 | Pause if assignment differs from the Active Register, unexpected users are affected, hard-stop behaviour is not observed, or rollback cannot be demonstrated. |
|---|---|
-

6.3 Apply scoped metered budgets

- | | |
|---|--|
| 1 | Confirm scope, US-dollar amount, period, covered_ai_credit_sku, owner, thresholds, alert_recipients, response_sla, Stop usage decision, and provider support. If an organization budget is used, record license-billing allocation evidence and the multi-organization caveat. |
| 2 | Apply the cost-center or organization configuration using an authorized administrator. Capture configuration snapshot and source timestamp. |
| 3 | Run P-04 in the approved controlled scope. Verify threshold notification and, if enabled, Stop usage outcome. |
| 4 | Reconcile provider state to the Active Register. Pause if a license baseline or cost-center included-usage control is being mistaken for a post-pool paid-usage guardrail or if evidence is incomplete. |
-

6.4 Apply Enterprise spending limit

- | | |
|---|--|
| 1 | Confirm enterprise billing entity, US-dollar amount, covered_ai_credit_sku, period, Enterprise budget owner approval, alert_recipients, response_sla, Stop usage intent, and emergency response. Record executive sponsor acceptance only for exceptional exposure or a rollback decision. |
| 2 | Apply through an authorized enterprise administrator. Capture configured state, source timestamp, approval, and evidence link. |
| 3 | Run the approved threshold and Stop usage test. Verify coordinated response with active scoped-budget controls. |
| 4 | Close the wave only after the Active Register, independent verification, support communication, and rollback owner agree. |
-

Rollout wave exit check

For every applied control: configured state, effective state, scope, unit, owner, test result, evidence link, exception status, and rollback reference match the Active Register.

PHASE PURPOSE

Operate sustains the approved control state through weekly triage, monthly reconciliation, analysis, feedback, training, and continuous governance. Monitor & Respond and Review Decision are recurring governance-loop components that support every lifecycle phase and are sustained through Operate.

WHAT THIS PHASE DELIVERS

- Weekly signals triaged with source, unit, scope, named owner, and response evidence.
- Monthly baseline reconciled with variance, source quality, decision owner, and due date.
- Analysis, adjustment, training, and improvement actions tied to evidence.
- Continuous Monitor & Respond and Review Decision records maintained in the Active Register.

Exit gate

The register, provider state, evidence, owners, monitoring signals, and review schedule remain current.

7.1 Weekly operational triage

Use weekly review for active alerts, rapidly changing burn rates, high-exposure cohorts, pilots, incidents, failed sources, and assignments due to expire. Validate source freshness, budget base, unit, affected population, owner, and response.

Weekly data is partial-period. Compare like periods and investigate material change before recommending a permanent budget or profile adjustment. Early AI initiatives can have higher anomaly risk and forecast variance; shorten the planning horizon and revisit funding assumptions more often while the evidence matures. [7]

7.2 Monthly reconciliation

Use monthly reconciliation to close the period across provider usage, included credits, license inventory, metered charges, budgets, cost centers, organizations, and assignments. Set a data cutoff, retain the original snapshot, capture restatements, and separate observed, calculated, and forecast values.

Do not declare the period reconciled while material sources remain incomplete. Record source gap, likely impact, owner, and due date.

7.3 Analyze, adjust, train, and improve

Convert reconciled data into an explicit decision. Assess concentration, organization mix, forecast variance, overage, alerts, exceptions, and user feedback. Account for population, licensing, rate, and feature changes before attributing outcomes to policy or training.

For a material initiative, record the intended business outcome and a proportionate value measure: strategic alignment, productivity or benefit result, compliance effectiveness, or cost-per-use when the organization has reliable telemetry. Do not require token- or inference-level measures when the chosen GitHub evidence sources do not expose them. [7]

Maturity overlay — not a replacement lifecycle

Use **Start** to establish baseline guardrails, owners, and showback; **Adopt** to add cost-center accountability, reviewed attribution, and chargeback where appropriate; and **Advance** to automate repeatable administration and forecasting after evidence is stable. Map this maturity overlay to the seven fixed phases in this guide; do not replace Prepare through Operate with a second sequence. [6]

Continuous governance loop

Monitor & Respond and Review Decision support every phase and are sustained through Operate

Monitor & Respond

Validate signals and take accountable action

Trigger	Required response
Example: 75% of named control base	Validate source, unit, scope, forecast, and owner; notify the accountable owner of the decision date.
Example: 90% of named control base	Verify assignments, current provider state, and support plan; prepare the tested response or rollback path.
Example: 100% of named control base	Determine whether the control is a ULB hard stop or an aggregate budget with Stop usage enabled; execute the approved business-continuity and support response.

Record every signal

Capture signal source and timestamp, base and unit, owner, response, provider state, user impact, decision, reconciliation result, alert_recipients, and response_sla in the Active Register.

Review Decision

Use reconciled consumption, outcomes, policy behaviour, exception history, and current provider capability to make one explicit decision. **Retain** when evidence confirms the approved state; **revise** when objective remains valid but the design must change; **roll back** when exit criteria fail, interruption is unexpected, provider behaviour differs materially, or evidence remains incomplete. Record the decision, rationale, owner, next review date, communication, and required register update; change active state only after provider state and evidence agree.

Appendix A: Design and assignment criteria

Criterion	Design question
Scope	Which enterprise, organization, cost center, user, SKU, and period does the control cover?
Unit	Is this AI Credits for a ULB, entitlement, or included-usage cap; USD for aggregate metered spend; or a forecast amount for a license baseline?
Owner	Who approves, operates, receives alert, supports users, and decides rollback?
Evidence	What source snapshot, test, configured-state, and effective-state record proves the control?
Rollback	What prior state is restored, by whom, and within what approved window?

Appendix B: Control decision sequence

Use this page during incident triage and review when a Copilot request or spend threshold requires a decision.

1	Identify the applicable ULB: individual, then cost center, then universal. If exhausted, support the user and verify assignment and reset timing.
2	If ULB has headroom, determine whether shared included credits remain. If they do, capture a provider evidence snapshot. If a cost-center included-usage cap applies, record whether the configured response is block or paid overage.
3	When paid overage is permitted after an included-usage cap, or when the shared pool is exhausted, identify cost-center, organization, and enterprise spending-limit scope. Verify US-dollar amount, owner, thresholds, allocation evidence, and Stop usage state on the applicable metered budget.
4	The named control owner executes the approved response for the documented control path. If scopes conflict or provider evidence is incomplete, the enterprise budget owner adjudicates; capture evidence and retain, revise, or roll back through the governance loop. [2]

Appendix C: Lifecycle and evidence checklist

Lifecycle component	Evidence	Exit criterion
Prepare	Scope, owners, source list, included-credit baseline.	Roles and evidence responsibilities accepted.
Baseline	Current provider state, gaps, exceptions, Active Register records.	Current state reconciled; unknowns named.
Design	Target controls, units, tests, rollback path.	Proposal is complete and internally consistent.
Approve	Authorization, risk decision, pilot scope.	Named approver and rollback owner recorded.
Pilot	P-01–P-05 results and effective-state evidence.	Pass, revise, or rollback decision documented.
Rollout	Wave record, configuration snapshot, verification.	Configured and effective state agree.
Operate	Signals, reconciliation, analysis, decision record.	Register and cadence remain current.
Continuous governance loop	Monitor & Respond and Review Decision records.	Signal, response, decision, owner, and evidence link are current.

Appendix D: Monitoring and response

During early experimentation or high forecast variance, use shorter forecast windows and revisit alert/response assumptions more often; do not treat a starter cadence or observed alert band as a universal provider default. [6] [7]

Define alert bases per control, not as a generic percentage. The supplied organization configuration shows 75%, 90%, and 100% thresholds as observed example values; confirm current supported threshold configuration in the provider UI and documentation for each aggregate metered budget. For ULBs, use internal monitoring or tenant-validated capability; do not describe provider percentage alerts as a sourced ULB feature without current evidence. Record the unit, named primary owner, delegate, executive escalation owner, on-call contact, source, evaluation period, notification path, canonical **alert_recipients**, canonical **response_sla**, and expected response. [2]

Exposure tier	Starter cadence
Low	Monthly operating review; quarterly design and ownership review.
Medium	Weekly operational review; monthly budget and profile review.
High	Twice-weekly operational review; monthly design, budget, and executive risk review.
Escalate tier when	Forecast uncertainty, active pilot, severe alert, failed evidence source, incident, or material exception occurs.

Appendix E: Reconciliation and data quality

- Preserve source system, retrieval method, source timestamp, freshness, calculation method, and known limitations.
- Retain an attribution map that links available usage and cost evidence to the applicable user, cost center, organization, and decision owner; state any mapping gap and its impact. [6] [7]
- Reconcile license inventory, entitlement pool, metered usage, provider adjustments, organization/cost-center assignment, and billing records.
- Maintain a restatement record: original snapshot, reason for change, revised value, owner, approval, and decision impact.
- Do not create a false final total from incomplete sources; retain unknowns and set an owner and due date.

Appendix F: Exceptions, change control, and rollback

1	Link the request to the Active Register record and current evidence baseline.
2	Capture prior provider state, rollback method, and expected restoration window.
3	Assess financial, security, entitlement, business, support, and interruption effects.
4	Obtain approvals, run controlled test where required, and apply through an authorized administrator.
5	Independently verify configured and effective state, observe for the approved period, and close only when evidence and register agree.

Appendix G: Review cadence

Cadence	Focus
Weekly	Active alerts, pilots, high-exposure cohorts, incidents, exceptions, and data freshness.
Monthly	Entitlements, usage, cost, forecast, assignments, reconciliation, and budget/profile decisions.
Quarterly	Profile design, access posture, ownership, evidence, training, and drift.
Event-driven	Rate, entitlement, feature, organization, risk, incident, or control change.

Appendix H: Getting-started checklist

1	Name approval, operating, support, escalation, and rollback decision owners.
2	Capture the included-credit baseline with timestamped provider evidence and the documented UTC reset rule.
3	Create and validate the Active Register template: entitlement-baseline, control, test, and exception record types; logical key; naming prefixes; dates; evidence links; and rollback references.
4	Assess current policy, budget, ULB, included-usage control, license baseline, cost-center, and assignment state.
5	Design financial layers, ULB taxonomy, policy profiles, included-usage controls, and metered budgets with units and owners.
6	Approve a bounded pilot with P-01 through P-05 success criteria.
7	Run the pilot, reconcile evidence, make a rollout decision, and then apply controlled waves.
8	Start the operating cadence and keep the Active Register current.

Appendix I: Governance glossary

Term	Definition
Active Register	The controlled system of record for typed policy-profile, ULB, entitlement-baseline, included-usage-control, metered-budget, license-baseline, rollout-wave, controlled-test, and exception records. It requires access control, change history, evidence links, and export capability.
Entitlement-baseline	Shared included-credit or entitlement evidence for one billing scope. It records entitlement basis, cycle/reset evidence, owner, source, and a forecast/reconciliation-only purpose; it is not an enforcement control.
Profile Register	The policy-profile filtered view of the Active Register; it is not a separate data model or logical key.
AI FinOps scope	A focused operating lens for AI spending that connects allocation, forecasting, optimization, governance, and business value. This guide applies that lens to GitHub Enterprise Cloud AI Credits; it does not represent every AI cost category. [7]
Showback	Publication of attributable consumption and cost information to the accountable owner for awareness and decision-making; it does not itself transfer an internal financial charge. [6] [7]
Chargeback	An approved internal allocation of attributable consumption or cost to a business owner or cost center. It requires a documented basis and finance process; it does not change GitHub provider control behavior. [6] [7]
Value measure	A proportionate, organization-defined outcome indicator such as strategic alignment, productivity or benefit result, compliance effectiveness, or cost-per-use where reliable telemetry exists. [7]
AI Credit	GitHub billing unit for eligible Copilot AI usage; GitHub documentation states one credit equals \$0.01 USD. [1]
Shared AI-credit pool	Monthly included-credit pool at the billing entity; shared rather than held in individual user buckets; resets at 00:00:00 UTC on the first day of each calendar month. [1]
User-Level Budget (ULB)	Per-user AI-credit boundary. Individual overrides cost center; cost center overrides universal. The applicable supported ULB is a hard stop. [2]
Cost-center ULB cohort	Per-user ULB for members of one cost center; it does not create a shared pool or extend the universal amount. [2]
Cost-center included-usage control	Provider-calculated cap on one cost center's draw from included AI Credits before the metered phase. At the cap, members are blocked or continue as paid overage; record the configured/observed behavior. It is distinct from a post-pool budget and has no aggregate-budget Stop usage state. [2] [5]
Cost-center metered budget	Aggregate USD spending guardrail for one cost center after shared-pool exhaustion. [2]
License-spend baseline	Forecast-only record of eligible licenses and expected license charges. It is not a ULB, included-usage cap, or paid-usage guardrail.
Organization metered budget	Aggregate USD guardrail for paid usage billed to one organization after pool exhaustion. Multi-organization licensing may make allocation unpredictable; preserve allocation evidence. [1] [2] [4]
Enterprise spending limit	Enterprise-wide aggregate USD control for paid AI-credit usage after pool exhaustion. Included credits are consumed first; only subsequent paid usage counts toward this control. [1] [2]
Stop usage	Optional setting, off by default, that can make a supported aggregate metered-spend budget a hard stop. [2]
Controlled test	Limited approved validation of scope, assignment, precedence, alerting, enforcement, and rollback.
Rollback	Restoration of captured prior provider state and active version after a failed exit gate.

Appendix J: Active Register schema and operating guide

The **Active Register** is the controlled system of record for every policy profile, User-Level Budget (ULB), entitlement-baseline, included-usage control, metered budget, license baseline, rollout wave, controlled test, and exception. It links the current provider configuration to the approved decision, the evidence proving effective behavior, the accountable people, and the rollback path. This appendix is complete on its own: it defines the schema, explains how to populate it, and gives the validation rules needed to operate it.

J.1 How the register works

- | | |
|---|--|
| 1 | Create one row when a control, pilot, rollout wave, or exception is proposed. Give it an internal <code>enterprise_control_id</code> before approval. |
| 2 | Populate baseline and design fields: scope, control type, unit, values, product coverage, owner, evidence source, and expected behavior. Use the canonical field dictionary in J.2. |
| 3 | At Approve, complete approval, effective-period, owner/delegate, support, and rollback fields. At Pilot, attach test results and configured/effective-state evidence. |
| 4 | At Rollout and Operate, update the same record with current provider state, alert or incident evidence, review decision, and a new version or exception record when the control changes. |
| 5 | Export a timestamped approved snapshot at each phase exit. The register is not a reporting convenience; it is the traceable decision record for the control lifecycle. |

Canonical identifier

The base logical key is (**record_type**, **scope_id**, **enterprise_control_id**). Append (**policy_or_profile**, **profile_version**) only when the record has a named profile/control pattern; otherwise use **not-applicable**, 0 in an export. **enterprise_control_id** is the organization's stable internal ID. Capture **provider_control_id** only when the provider exposes a native ID; otherwise capture the configuration/source URL and retrieval timestamp. It is evidence, not a substitute for the internal key.

J.2 Canonical field dictionary — identity and ownership

Field	Requirement / allowed value	How it works
record_type	Required for every row. Picklist: policy-profile, ULB, entitlement-baseline, included-usage-control, metered-budget, license-baseline, rollout-wave, controlled-test, exception.	Classifies record behavior and required evidence.
enterprise_control_id	Required for every row. Immutable internal identifier, for example ENT-CTRL-0003.	Links approvals, tests, deployment, and review without relying on mutable provider labels.
provider_control_id	Required only when the provider exposes a native ID. Otherwise leave unavailable and capture source_url or configuration URL plus retrieval_timestamp.	Connects the record to provider configuration evidence without becoming a logical-key field.
scope_id	Required immutable enterprise, organization, cost-center, or user identifier; never a display name alone.	Defines the controlled population or billing scope.
assignment_target	Required when a control is assigned to users, organizations, or rollout waves.	Names the effective target without changing the scope ID.
policy_or_profile	Conditional: required for policy-profile; use when a named control/profile pattern exists; otherwise not-applicable.	Names the internal policy/profile or guardrail pattern without becoming a universal requirement.
profile_version	Conditional: required for policy-profile and named ULB/control patterns; otherwise use 0 in an export key suffix. Increase for a material pattern change.	Enables approval and rollback to a prior named design without applying to unrelated record types.
owner_primary	Required from Prepare. Named person accountable for the control lifecycle.	Owens the record and its scheduled review.
owner_delegate	Required from Approve for production-intended records. Named backup for alerts and decisions.	Prevents a control from becoming ownerless.
support_contact	Required from Approve for production-intended records.	Receives user-impact and interruption escalation.

J.3 Canonical field dictionary — controls, timing, and evidence

Field	Requirement / allowed value	How it works
ulb_type	Required only for ULB: universal, cost-center, or individual.	Shows the applicable per-user scope and supports precedence review.
ulb_amount_ai_credits	Required only for ULB; whole AI-Credit amount.	The value used for the ULB hard-stop test.
budget_amount_currency	Required only for metered-budget; numeric-only decimal with no currency symbol or text.	The aggregate post-pool budget amount; use a separate currency field.
entitlement_basis / included_credits_amount / baseline_purpose	Required only for entitlement-baseline. Entitlement basis is the provider licensing or billing basis; included-credit amount is recorded when exposed; baseline_purpose is forecast-only, reconciliation-only, or both.	Makes the shared included-credit baseline a traceable, non-enforcement record.
currency	Required for metered-budget rows; value is USD.	GitHub documents additional-usage budgets in USD; one AI Credit equals \$0.01 USD. [1]
license_baseline_amount	Required only for license-baseline; numeric forecast amount without currency text.	Forecast-only input; never an enforcement boundary.
license_baseline_currency	Required only for license-baseline; ISO invoice-currency code, for example USD.	Stores the currency for license_baseline_amount. It is separate from metered-budget currency.

J.3 continued — coverage, alerting, timing, and evidence

Field	Requirement / allowed value	How it works
covered_ai_credit_sku	Required for metered-budget and recommended for ULB. One or more enumerated provider SKU labels or IDs; encode multiple values as a semicolon-separated controlled list.	Record a dated coverage decision. Store umbrella UI selectors as observed evidence or notes, not as generalized canonical values.
alert_thresholds	Required for an alerting aggregate control; ordered semicolon-separated provider-supported values with units, for example 75%;90%;100%.	Defines alert bands against the named control base. The example values are observed configuration evidence, not universal provider defaults.
alert_recipients	Required for an alerting metered-budget; semicolon-separated named roles, groups, or validated addresses.	Stores the accountable recipient set observed or configured for the threshold notification.
response_sla	Required for an alerting metered-budget; approved response target such as within 1 business day.	States the expected response time and escalation expectation for the alert owner.
stop_usage_state	Required only for metered-budget: enabled, disabled, unavailable, or not-verified.	Records the separate aggregate-budget hard-stop intent; never assign it to an included-usage-control.
downstream_metered_budget_reference (display: Downstream metered-budget reference)	Optional only when an included-usage-control permits paid overage.	Links the paid-overage route without implying that the included-usage control has an aggregate-budget Stop usage setting.
billing_cycle_start / billing_cycle_end / observed_next_reset / timezone_source	ULB: record observed_next_reset plus source/timezone when exposed; do not infer cadence. Included-usage-control: record cycle/reset evidence for the selected scope. Entitlement-baseline: record billing_cycle_start, billing_cycle_end, observed_next_reset, and timezone_source when shown.	Preserves evidence by record type. The published shared-pool monthly reset does not establish ULB reset timing.
approval_id / effective_start / effective_end	Required from Approve for production-intended records.	Defines the approved change and its active period.
test_id / test_result	Required from Pilot: P-01, P-01a, P-02, P-03, P-04, or P-05; result is pass, fail, blocked, or not-run.	Links the record to observable validation.
evidence_link / source_url / retrieval_timestamp	Required from Baseline. Use a durable URL, repository path, or evidence-system ID and ISO 8601 timestamp.	Allows an independent reviewer to reproduce the evidence.
rollback_reference	Required from Approve for production-intended records.	Points to the prior state and approved restoration action.
exception_id / exception_expiry / record_status / notes	Use when an approved exception, retirement, gap, or limitation exists.	Makes deviation visible and time-bounded.

Canonical field token

downstream_metered_budget_reference is the exact export and CSV field name. Use it only when an included-usage control permits paid overage and a separate metered-budget record governs that paid route. A block-at-cap configuration has no downstream paid route. The token contains no spaces or hyphens.

J.4 Record-type requirements and validation rules

Record type	Mandatory content	Validation rule
policy-profile	enterprise_control_id; provider_control_id when exposed, otherwise source_url + retrieval_timestamp; scope_id; policy_or_profile; profile_version; owners; approval; P-01a; evidence; rollback.	Policy/profile name and version cannot change after approval; create a new record for a material revision.
ULB	enterprise_control_id; scope_id; assignment_target; ulb_type; ulb_amount_ai_credits; profile_version; cycle/reset evidence; P-01 and P-03; owners; rollback.	ULB amount is an integer in AI Credits. Verify coverage per intended product surface; a cost-center cohort must have a validated membership report.
entitlement-baseline	enterprise_control_id; provider_control_id when exposed, otherwise source_url + retrieval_timestamp; billing_scope; entitlement_basis; included_credits_amount when exposed; billing-cycle start/end; observed_next_reset; timezone_source; owner; evidence; baseline_purpose.	Forecast and reconciliation only. It is not a ULB, included-usage cap, or paid-usage guardrail; baseline_purpose must be forecast-only, reconciliation-only, or both.
included-usage-control	enterprise_control_id; provider_control_id when exposed, otherwise source_url + retrieval_timestamp; scope_id; directly assigned license evidence; provider-calculated AI-credit cap; configured/observed behavior at cap (block or paid-overage); optional downstream_metered_budget_reference only for paid-overage; P-02; owners; evidence.	It is pre-metered and provider-calculated. It has no aggregate-budget Stop usage setting. At the cap, members are blocked or continue as paid overage; only paid overage may use a referenced separate metered-budget record.
metered-budget	enterprise_control_id; provider_control_id when exposed, otherwise source_url + retrieval_timestamp; scope_id; budget_amount_currency numeric-only; currency=USD; covered_ai_credit_sku (semicolon-separated labels/IDs when multiple); alert_thresholds; alert_recipients; response_sla; stop_usage_state; P-04; owners; rollback.	Do not mark Stop usage enabled until a current provider evidence link and test result are present. Reject non-numeric budget_amount_currency values such as "20 USD".
license-baseline	enterprise_control_id; provider_control_id when exposed, otherwise source_url + retrieval_timestamp; scope_id; covered license SKU; eligible licenses; license_baseline_amount numeric-only; license_baseline_currency; period; variance owner; evidence.	Forecast-only. It cannot be used as a ULB, included-usage cap, or paid-usage hard-stop.
rollout-wave	enterprise_control_id; scope_id; assignment_target; approved profile or control version; effective period; communication; verification; rollback.	Close only after configured and effective state agree with the Active Register.
controlled-test	enterprise_control_id; test_id; pilot population; period; configured/effective state; evidence; result; decision owner.	No pass result without timestamped evidence and a defined blast radius.
exception	enterprise_control_id; affected control; scope_id; reason; approver; compensating control; expiry; review date.	An expired exception cannot remain active; renew or close it through a recorded decision.

J.5 Worked examples and operating pattern

This PDF is self-contained: Appendix J is the Active Register implementation definition; the CSV is optional convenience only. Protect canonical-key and identifier columns, use controlled picklists, validate numeric-only budget amounts and required evidence, restrict editors, retain history, and export a signed snapshot at each exit gate. Apply stop_usage_state only to metered-budget records. An included-usage-control records block or paid-overage behavior; use downstream_metered_budget_reference only for a paid-overage route. The example below is an enforceable cost-center metered budget. Entitlement-baseline and license-baseline records use the same base key; their type-specific fields and validation rules are defined in J.3/J.4. Values are illustrative; verify provider state before use.

Canonical field	Example
record_type	metered-budget
enterprise_control_id	ENT-CTRL-0003
provider_control_id	budget-ccr-0012-it-ai-credit-v1
scope_id	ccr-0012-it
policy_or_profile	ai-credit-guardrail
profile_version	v1
budget_amount_currency	40 (numeric-only)
currency	USD
covered_ai_credit_sku	Copilot AI credits; add further enumerated provider labels/IDs with semicolons when applicable; observed UI selector retained in notes
alert_thresholds	75%;90%;100% — observed configuration example; confirm provider-supported values
alert_recipients	cost-center-owner; finance-delegate
response_sla	within 1 business day for threshold response
stop_usage_state	enabled after P-04 evidence
owner_primary	cost-center-owner
owner_delegate	finance-delegate
test_id	P-04
test_result	pass
rollback_reference	CHG-1042; prior-state snapshot 2026-08-25
evidence_link	https://records.example/evidence/ccr-0012-it/2026-08-26

Operating pattern. At phase start, filter the register to the applicable record types and scope and confirm the prior exit gate. At phase exit, validate mandatory type fields, attach evidence, record the owner/delegate and review date, then export the signed phase snapshot. When a setting, amount, scope, coverage decision, or behavior changes materially, retain the prior row, create the revised version, repeat approval and pilot as required, and link the rollback reference.

References

Use current GitHub documentation as the source of truth for product capabilities, billing, support, and administrative role requirements. Use the GitHub Well-Architected guidance and FinOps Foundation framework as cited design and operating guidance; neither replaces current provider documentation.

	Source
[1]	GitHub Docs. <i>Usage-based billing for organizations and enterprises</i> . https://docs.github.com/en/enterprise-cloud@latest/copilot/concepts/billing/usage-based-billing-for-organizations-and-enterprises
[2]	GitHub Docs. <i>Budgets for usage-based billing</i> . https://docs.github.com/en/copilot/concepts/billing/budgets-for-usage-based-billing
[3]	GitHub Docs. <i>Getting started with budget controls</i> . https://docs.github.com/en/copilot/tutorials/budgets/getting-started-with-budget-controls
[4]	GitHub Docs. <i>Optimizing your budget configuration</i> . https://docs.github.com/en/copilot/tutorials/budgets/optimizing-your-budget-configuration
[5]	GitHub Docs. <i>Cost centers</i> . https://docs.github.com/en/billing/concepts/cost-centers
[6]	GitHub Well-Architected. <i>Managing AI credits</i> . Updated September 14, 2026. https://learn.github.com/well-architected/governance/recommendations/managing-ai-credits
[7]	FinOps Foundation. <i>FinOps for AI</i> . FinOps Framework technology category. https://www.finops.org/framework/technology-categories/ai/

Reference note

Accessed September 15, 2026. The named organizations and configuration screenshots are illustrative operational examples. GitHub Well-Architected and FinOps Foundation material is incorporated as design and operating guidance. Provider capabilities, UI labels, pricing, and administrative roles may change; validate them in current GitHub documentation before production use.